

Minimax Decision

function Minimax-Decision (game) **returns** operator

return operator that leads to a child state with the
 max Minimax-Value(child state,game))

function Minimax-Value(state,game) **returns** utility value

if Goal(state), **return** Utility(state)
else if Max's move then
 → **return** max of successors' Minimax-Value
else
 → **return** min of successors' Minimax-Value



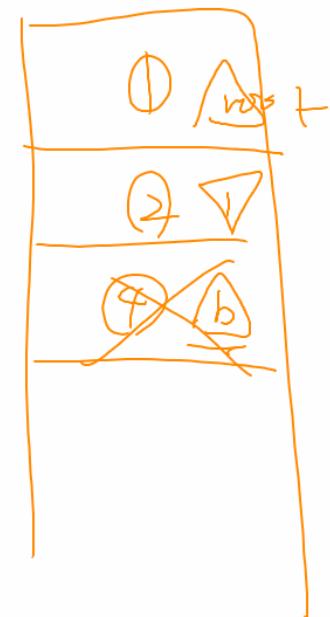
Minimax Decision

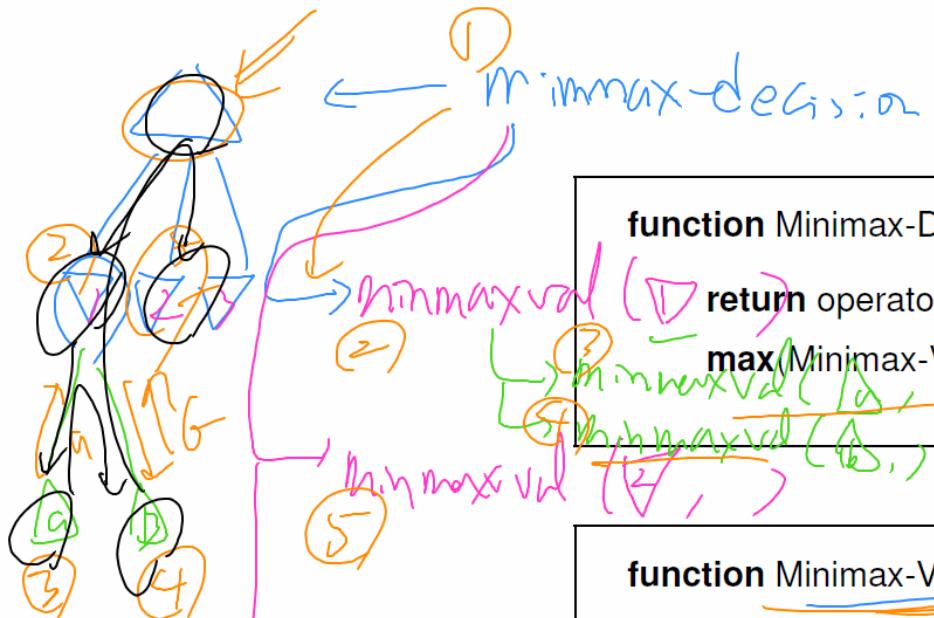
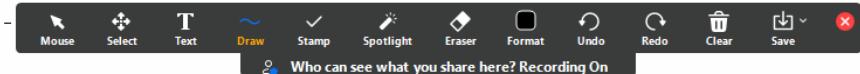
function Minimax-Decision (game) **returns** operator

return operator that leads to a child state with the
 max Minimax-Value(child state, game))

function Minimax-Value(state, game) **returns** utility value

if Goal(state), **return** Utility(state)
else if Max's move then
 → **return** max of successors' Minimax-Value
else
 → **return** min of successors' Minimax-Value





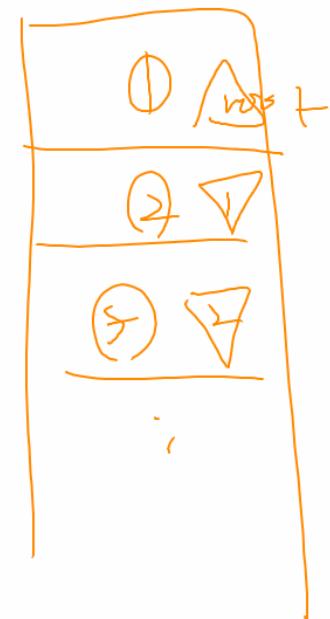
Minimax Decision

function Minimax-Decision (game) **returns** operator

return operator that leads to a child state with the
 max Minimax-Value(child state,game))

function Minimax-Value(state,game) **returns** utility value

if Goal(state), **return** Utility(state)
else if Max's move then
 → **return** max of successors' Minimax-Value
else
 → **return** min of successors' Minimax-Value



So here and then to here, then to here, so that's basically a depth, first exploration due to the way that this whole recursive thing is implemented

You are screen sharing Stop Share

Minimax Decision

```
function Minimax-Decision (game) returns operator
```

return operator that leads to a child state with the
 $\max(\text{Minimax-Value}(\text{child state}, \text{game}))$

```
function Minimax-Value(state,game) returns utility value
```

if Goal(state), **return** Utility(state)

else if Max's move then

→ **return** max of successors' Minimax-Value

else

→ **return** min of successors' Minimax-Value



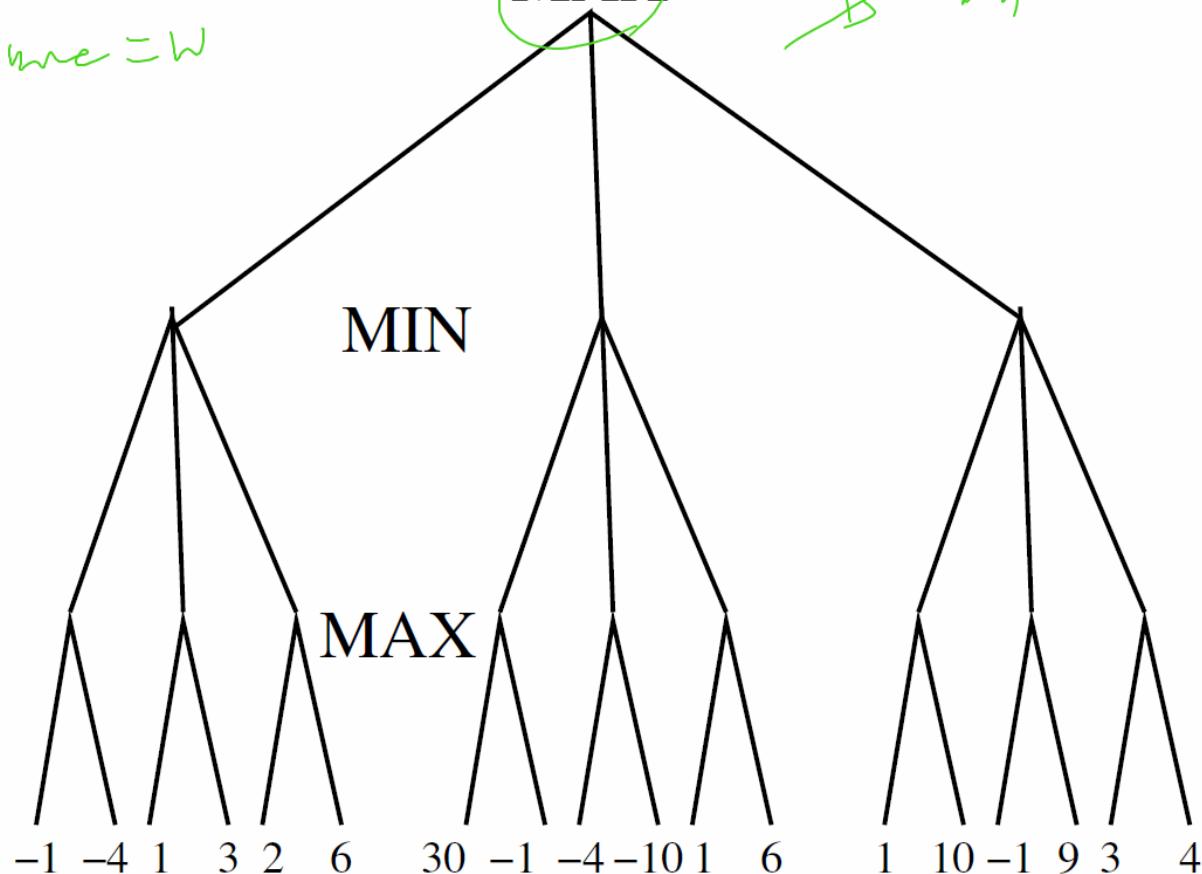


F. Max = B
F. min = W

Minimax Exercise

MAX

~~B~~ ~~W~~



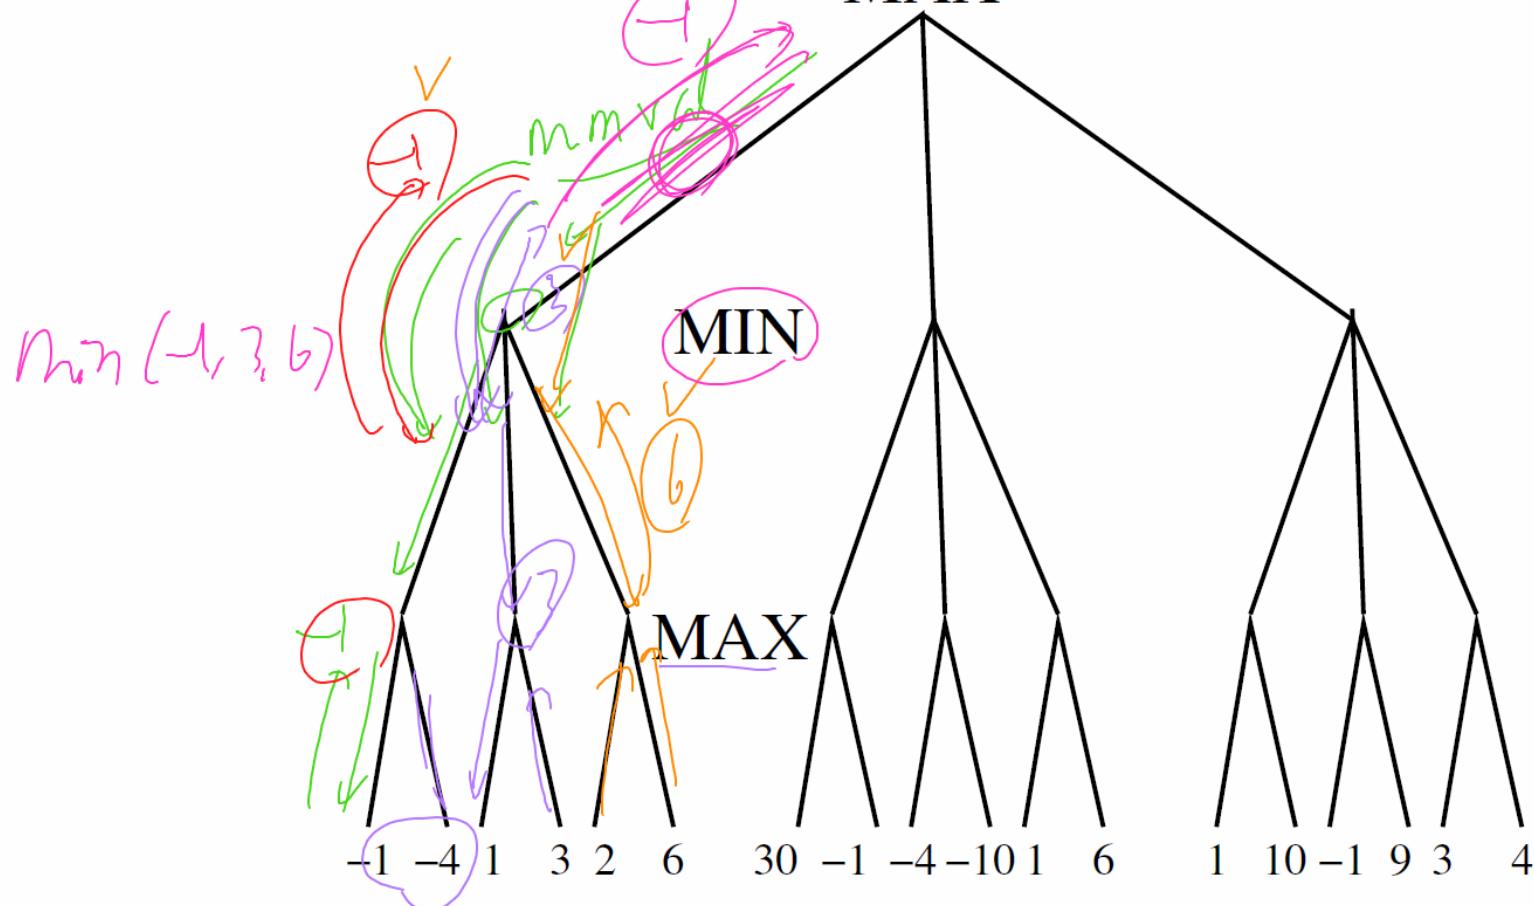
You always say the first place, wherever is the first? Yes.

You are screen sharing Stop Share



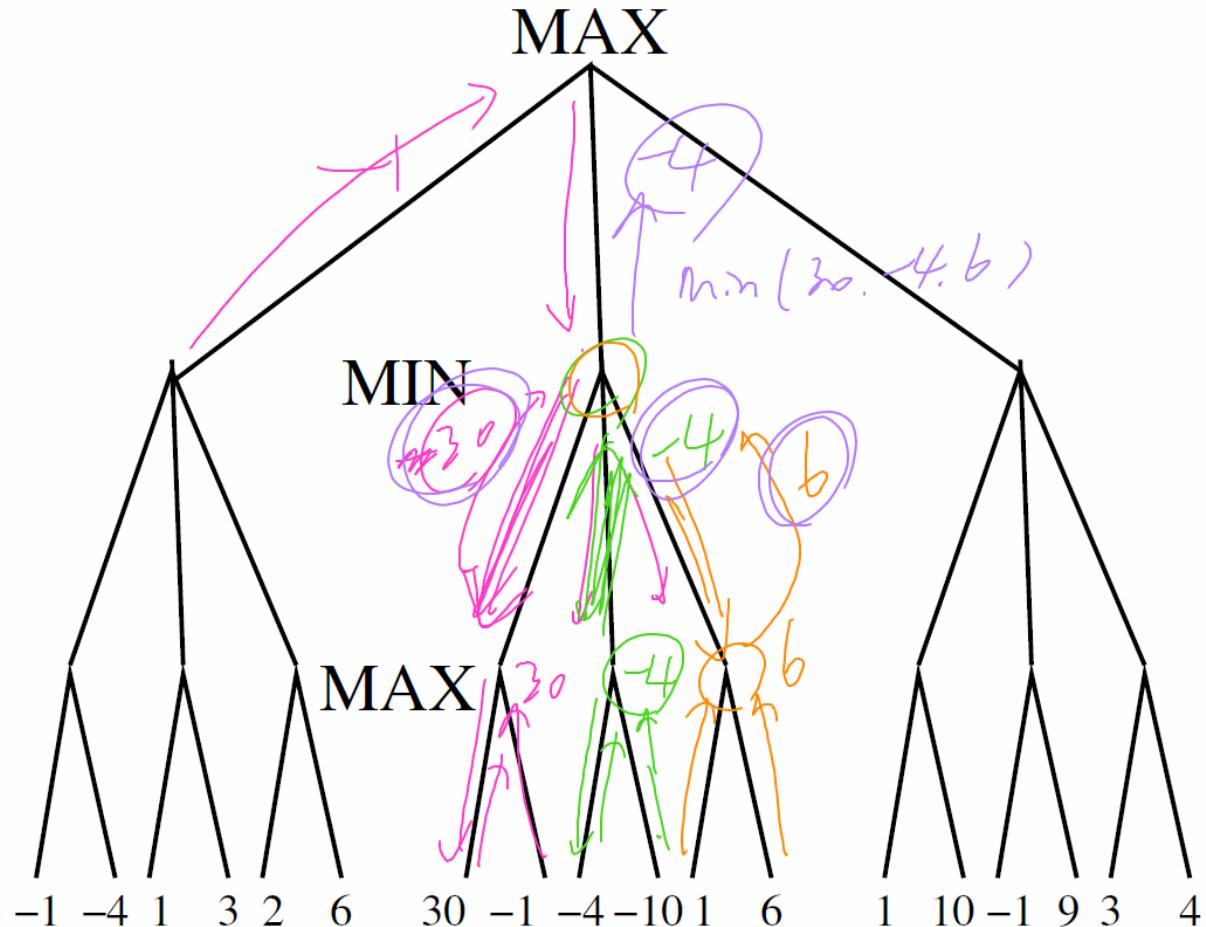
Minimax Exercise

MAX

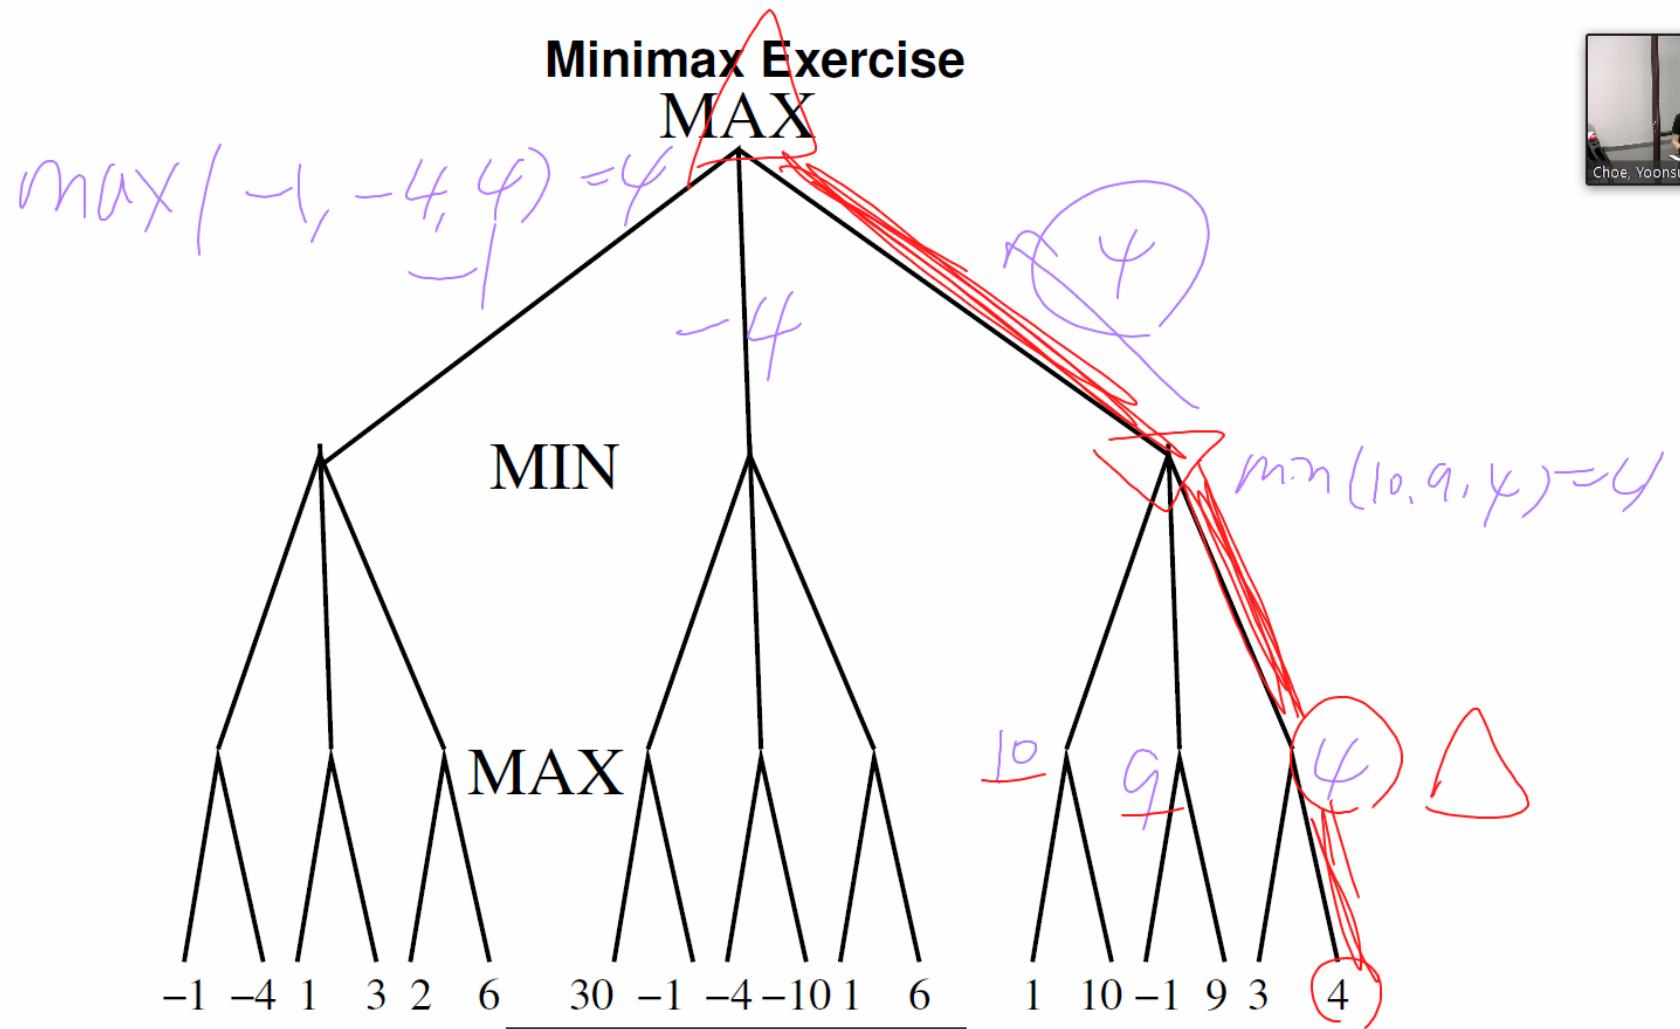




Minimax Exercise



Oh, 30, as 4 and 6 b minus 4,

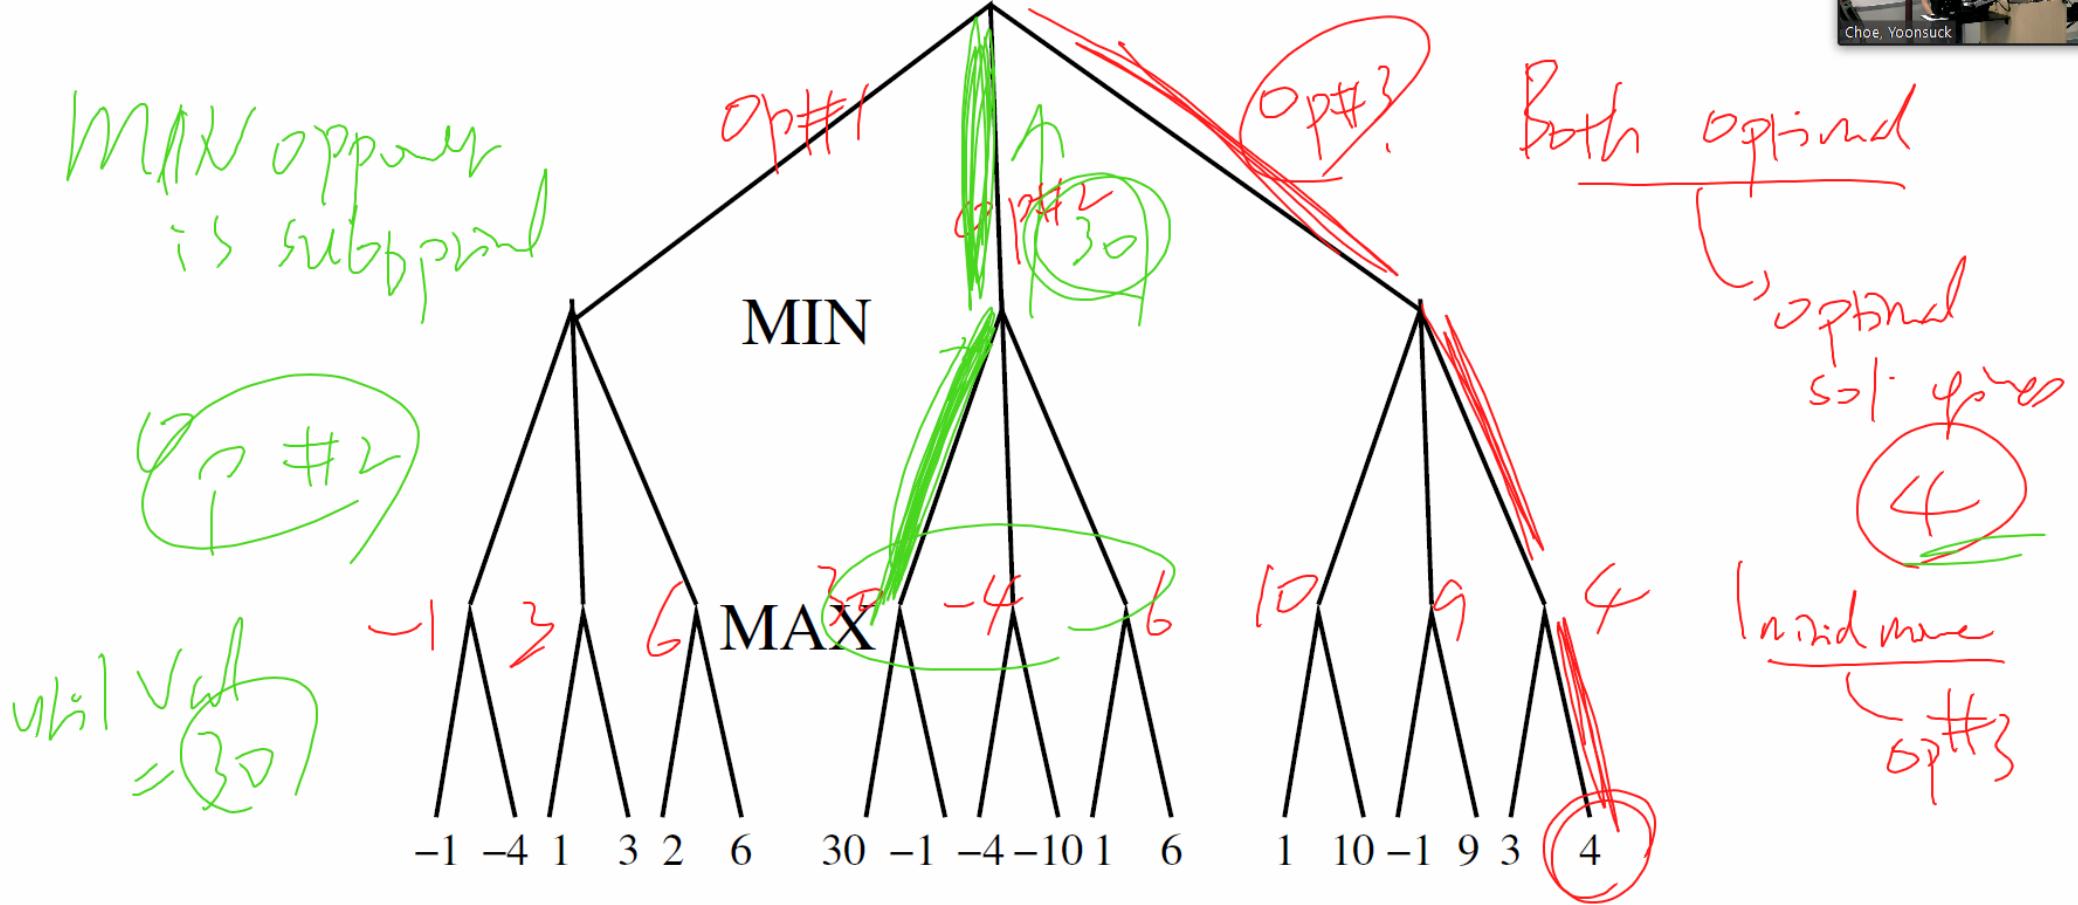


You'll pick 4, so that's the solution using this very simple recursive



Minimax Exercise

MAX





Resource Limits

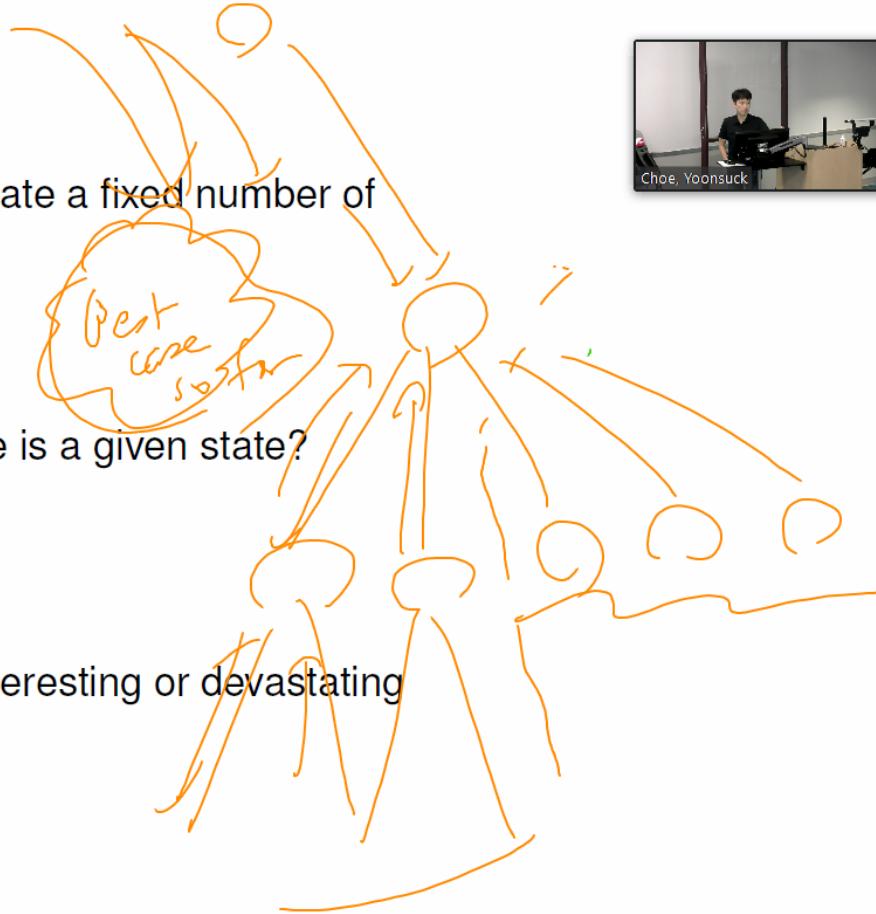
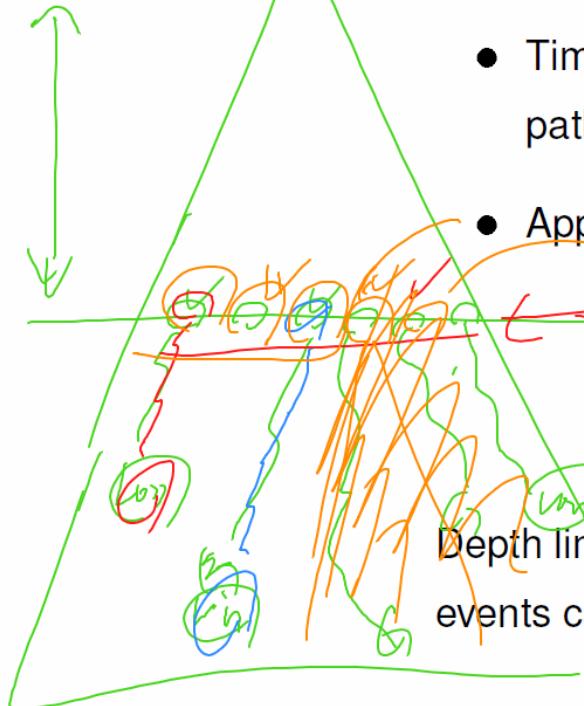
- Time limit: as in Chess → can only evaluate a fixed number of paths

- Approaches:

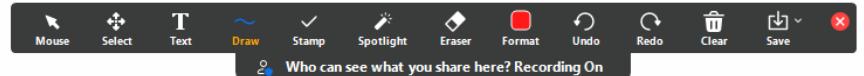
evaluation function : how desirable is a given state?

- **cutoff test** : depth limit
- **pruning**

Depth limit can result in the **horizon effect**: interesting or devastating events can be just over the horizon!



And suppose that a certain point you decide that whatever you do it's not going to be the the best case scenario so far, which is information that



Minimax Decision

```
function Minimax-Decision (game) returns operator  
    return operator that leads to a child state with the  
    max(Minimax-Value(child state,game))
```

```
function Minimax-Value(state,game) returns utility value  
    if Goal(state), return Utility(state)  
    else if Max's move then  
        → return max of successors' Minimax-Value  
    else  
        → return min of successors' Minimax-Value
```

Resource Limits



- Time limit: as in Chess → can only evaluate a fixed number of paths

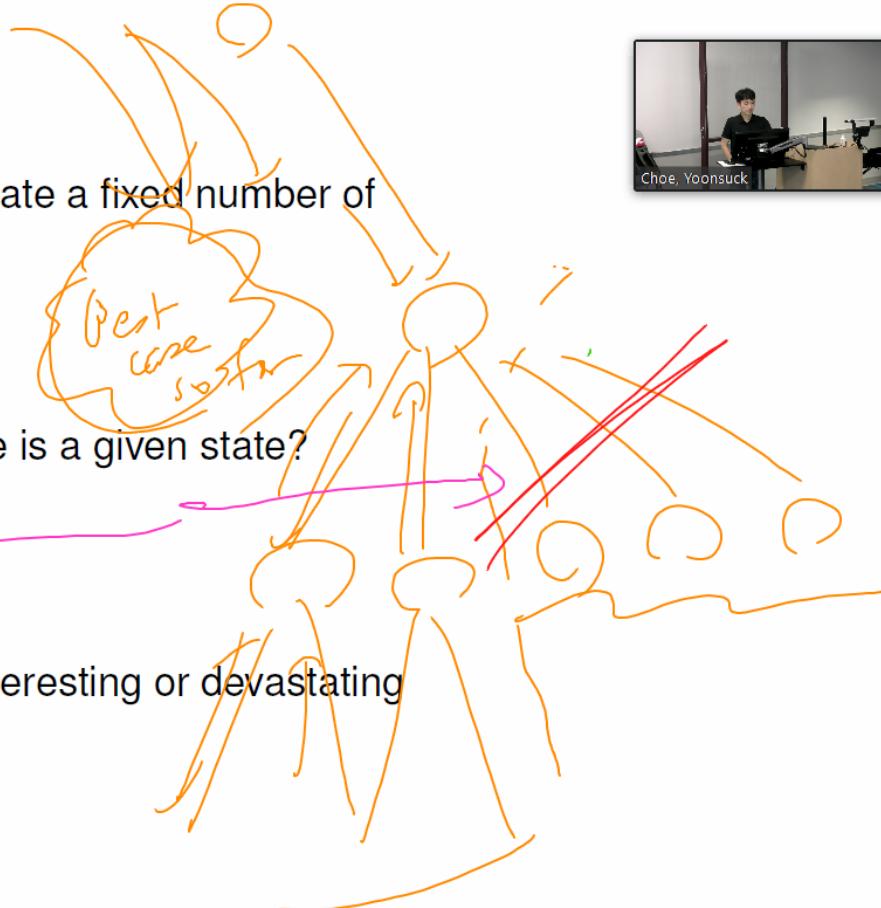
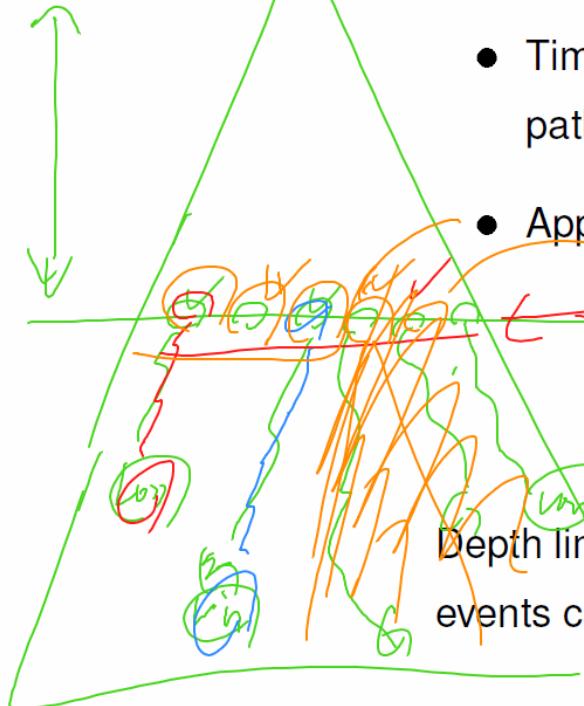
- Approaches:

evaluation function : how desirable is a given state?

- **cutoff test** : depth limit

- **pruning**

Depth limit can result in the **horizon effect**: interesting or devastating events can be just over the horizon!

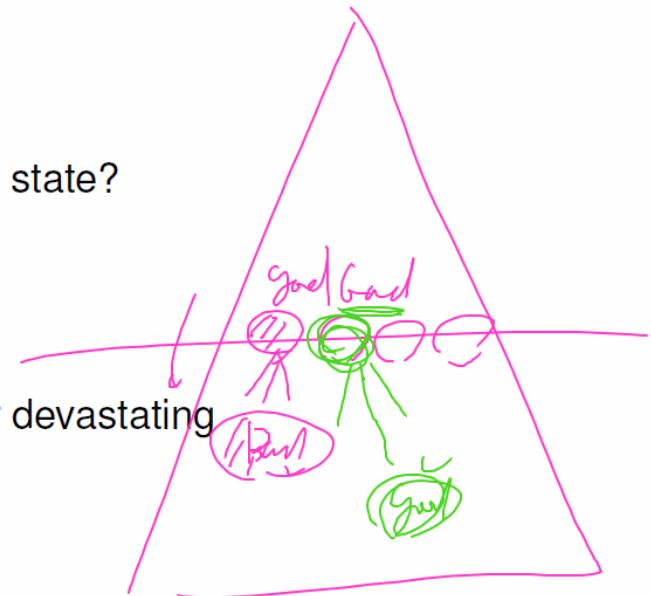


Resource Limits



- Time limit: as in Chess → can only evaluate a fixed number of paths
- Approaches:
 - **evaluation function** : how desirable is a given state?
 - **cutoff test** : depth limit
 - **pruning**

Depth limit can result in the **horizon effect**: interesting or devastating events can be just over the horizon!





Evaluation Functions

For chess, usually a **linear** weighted sum of feature values:

- $\text{Eval}(s) = \sum_i w_i f_i(s)$
- $f_i(s) = (\text{number of white piece } X) - (\text{number of black piece } X)$
- other features: degree of control over the center area *Rum*
- exact values do not matter: the **order** of Minimax-Value of the successors matter.

Evaluation Functions



For chess, usually a **linear** weighted sum of feature values:

- $\text{Eval}(s) = \sum_i w_i f_i(s)$
- $f_i(s) = (\text{number of white piece } X) - (\text{number of black piece } X)$
- other features: degree of control over the center area *Pawn*
- exact values do not matter: the **order** of Minimax-Value of the successors matter.

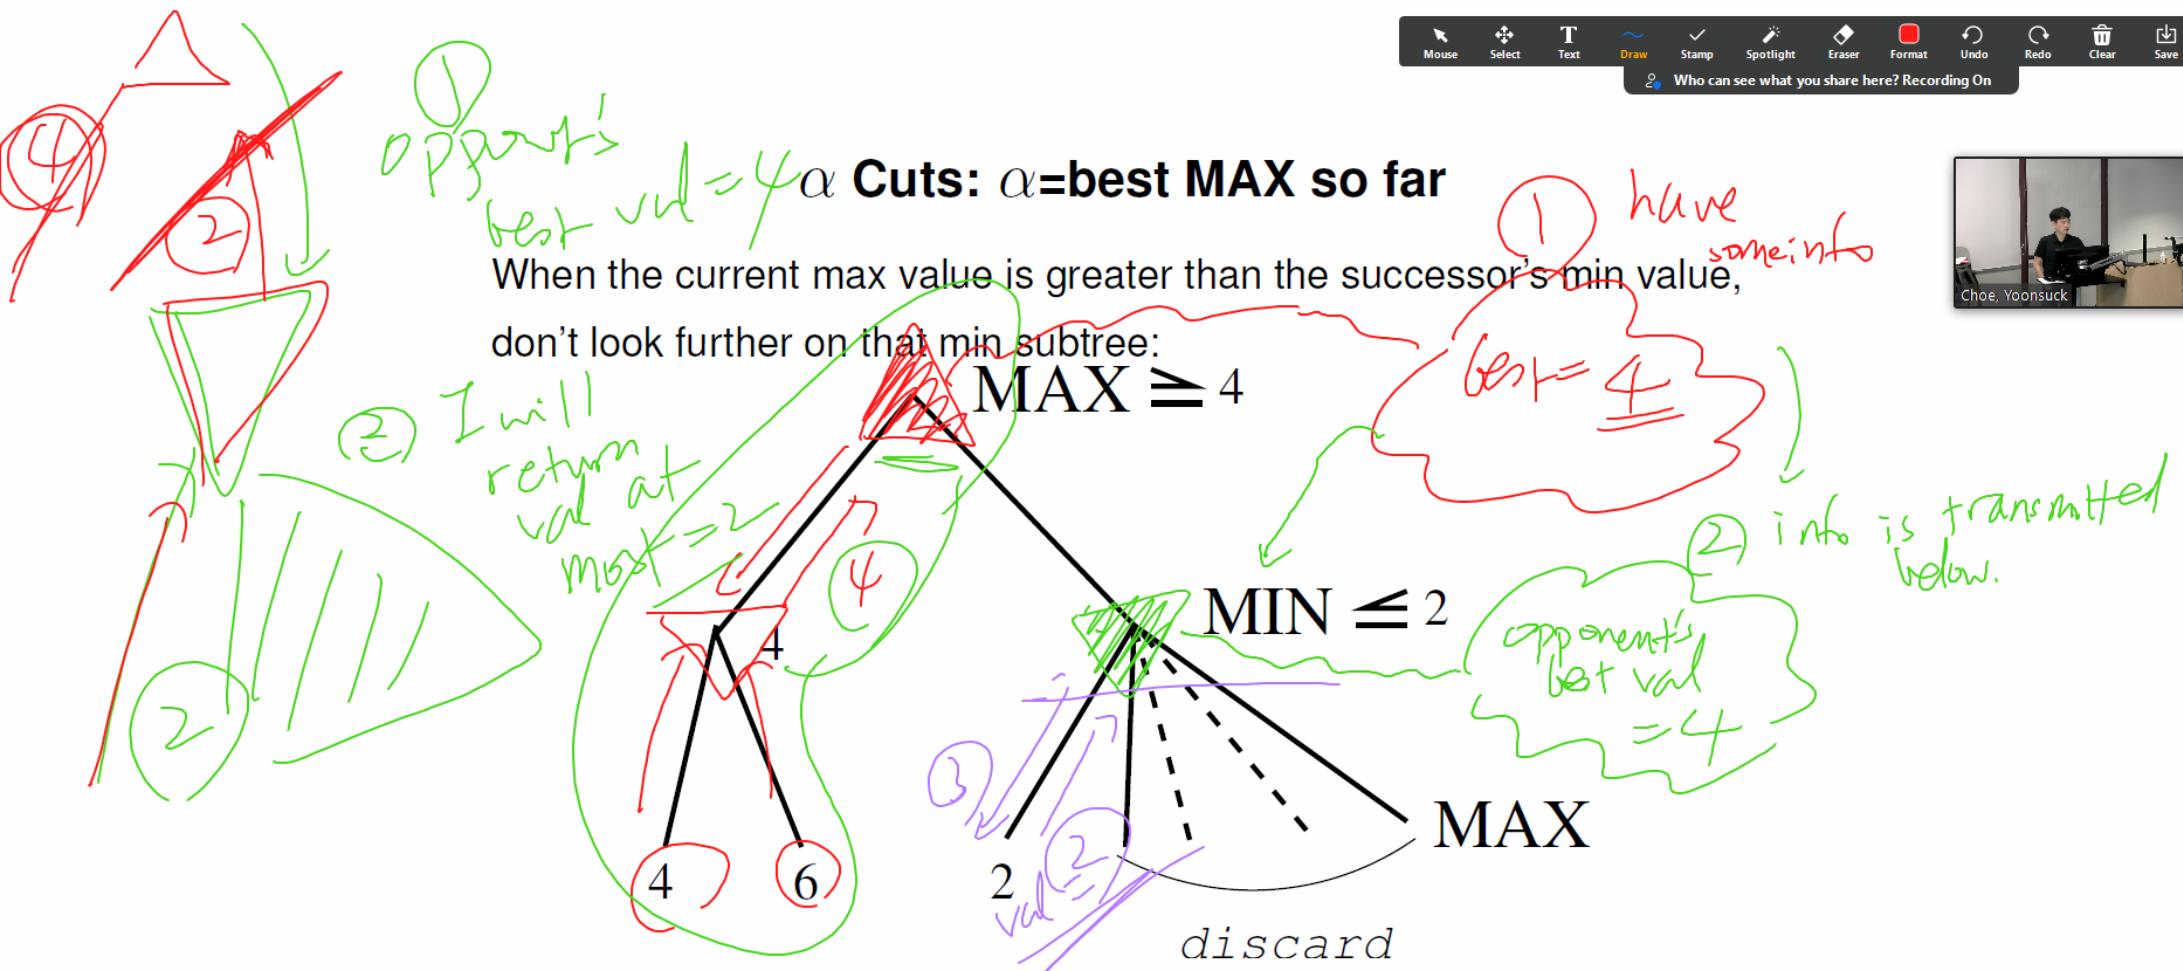
Evaluation Functions



For chess, usually a **linear** weighted sum of feature values:

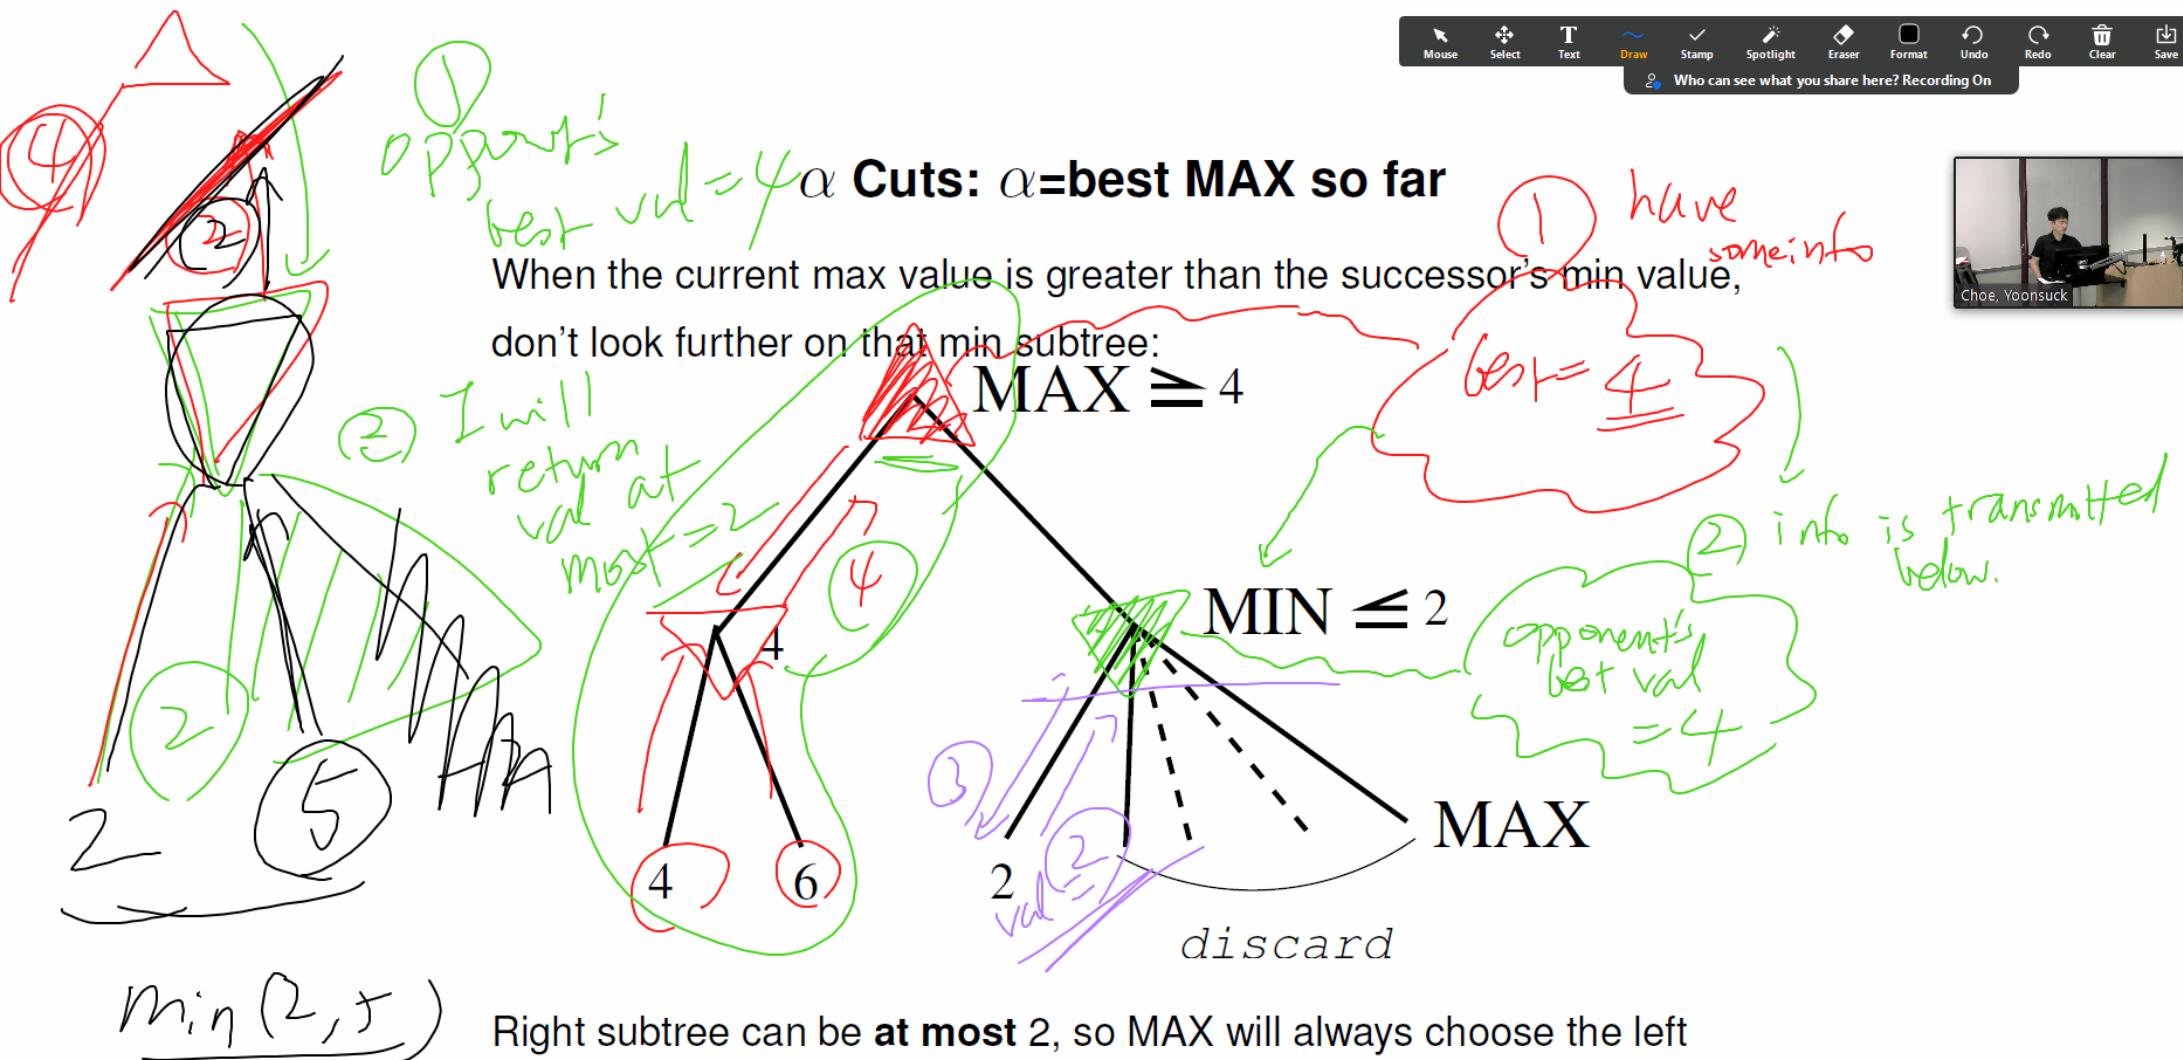
- $\text{Eval}(s) = \sum_i w_i f_i(s)$ $= 10, 20, 30 \checkmark$
 $= 5, 200, 3000 \checkmark$
- $f_i(s) = (\text{number of white piece } X) - (\text{number of black piece } X)$
- other features: degree of control over the center area
- exact values do not matter: the **order** of Minimax-Value of the successors matter.

so you lower it. If the ordering is the same, then for this particular case



Right subtree can be **at most 2**, so MAX will always choose the left path regardless of what appears next.

Yeah is so actually if you actually switch it's not better So if you return to, then, of course, your opponent will just ignore you, and just go for full





α Cuts: $\alpha = \text{best MAX so far}$

When the current max value is greater than the successor's min value, don't look further on that min subtree:

MAX ≥ 4

$\min(2, -8)$

4
4 6

MIN ≤ 2

discard

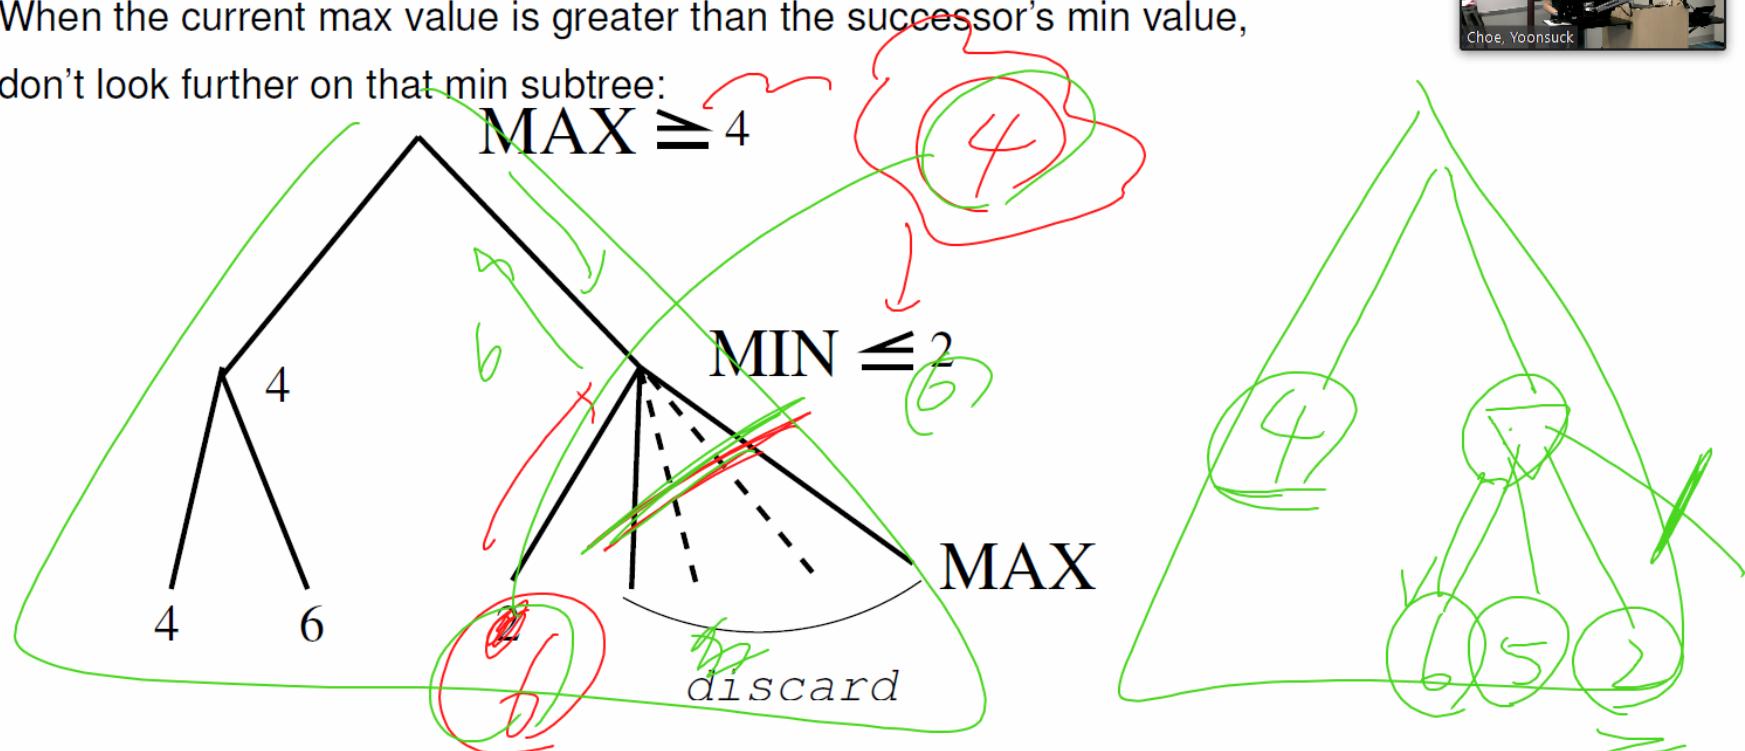
MAX

Right subtree can be **at most 2**, so MAX will always choose the left path regardless of what appears next.

α Cuts: $\alpha = \text{best MAX so far}$



When the current max value is greater than the successor's min value,
don't look further on that min subtree:



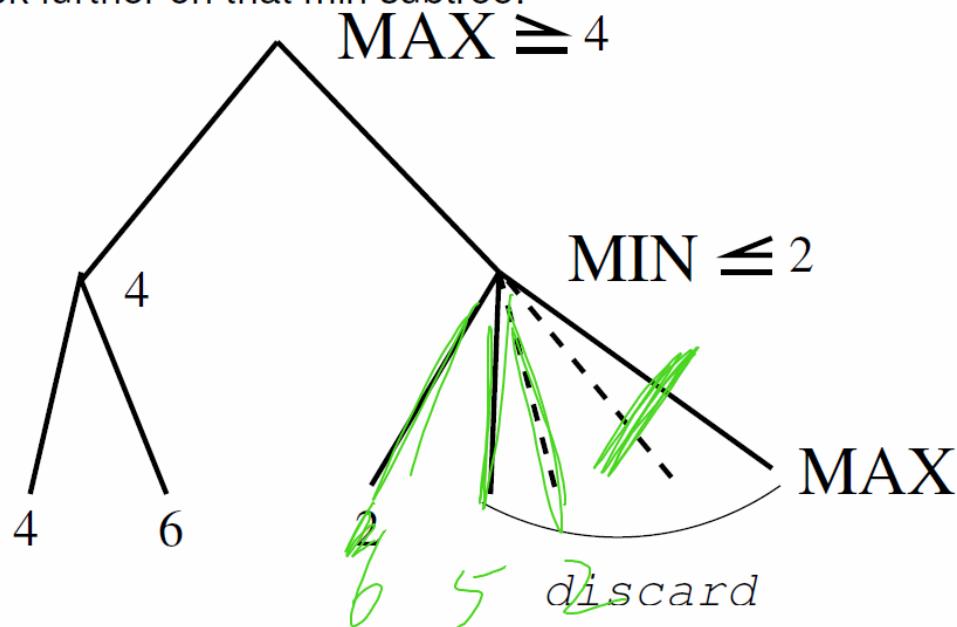
Right subtree can be **at most 2**, so MAX will always choose the left path regardless of what appears next.

Then you have to keep on going. Then, as soon as you reach that which is worse than that, then you can cut



α Cuts: $\alpha = \text{best MAX so far}$

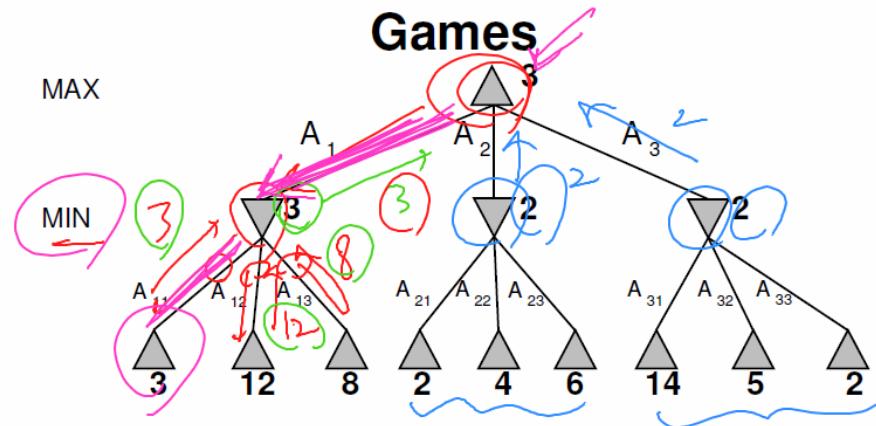
When the current max value is greater than the successor's min value, don't look further on that min subtree:



Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.

So to summarize. If this was 6 and 5, and then you have 2, then only after absorbing these 3, then you can cut

Minimax: Strategy for Two-Person Perfect Info Games



- generate the whole tree, and apply util function to the leaves
- go back upward assigning utility value to each node
- at MIN node, assign **min(successors' utility)**
- at MAX node, assign **max(successors' utility)**
- **assumption:** the opponent acts optimally

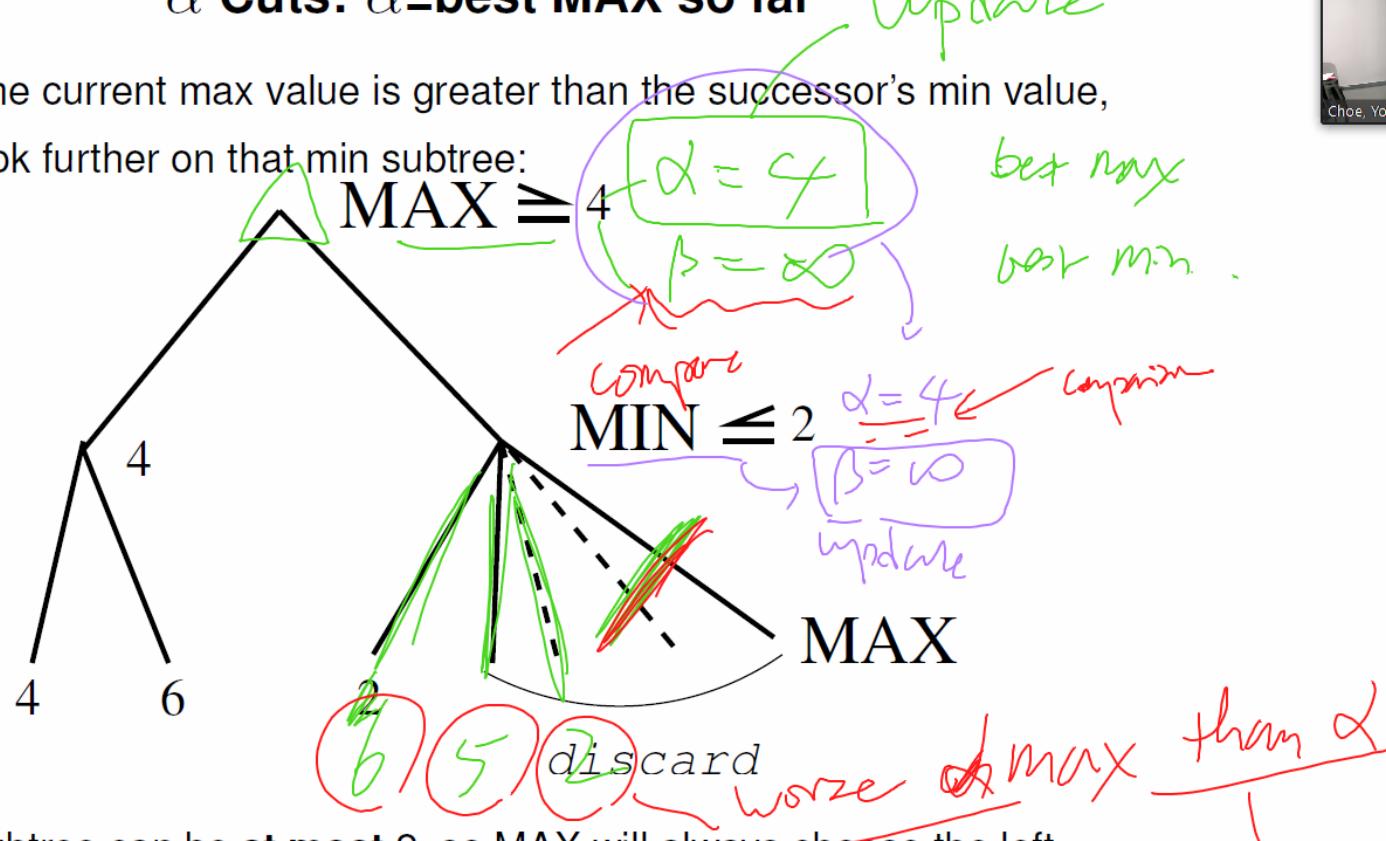


$$\alpha = -\infty$$

α Cuts: $\alpha = \text{best MAX so far}$

$$\beta = \infty$$

When the current max value is greater than the successor's min value,
don't look further on that min subtree:



Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.

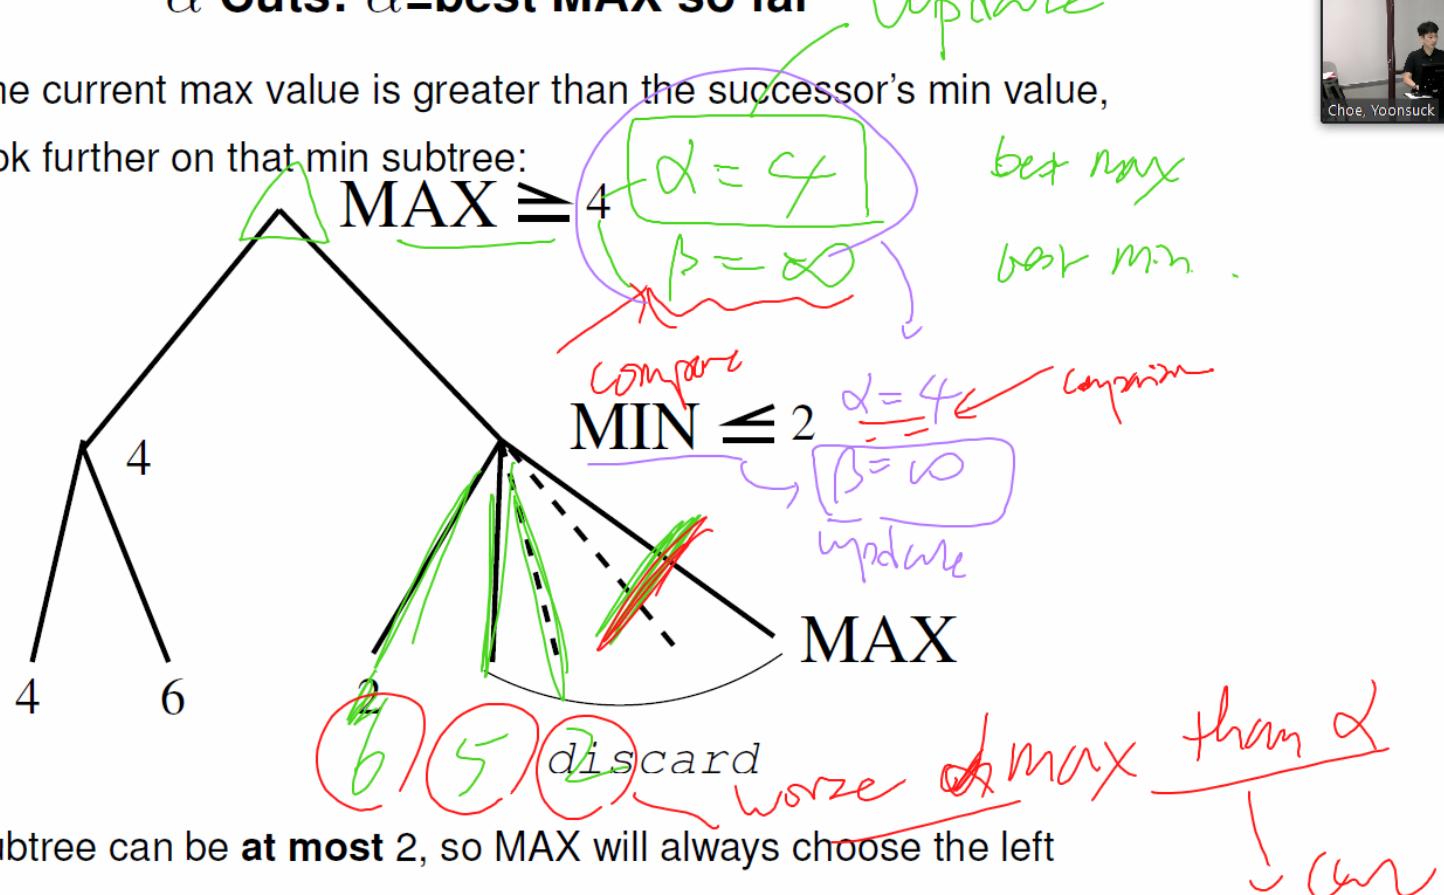


$$\alpha = -\infty$$

$$\beta = \infty$$

α Cuts: $\alpha = \text{best MAX so far}$

When the current max value is greater than the successor's min value,
don't look further on that min subtree:

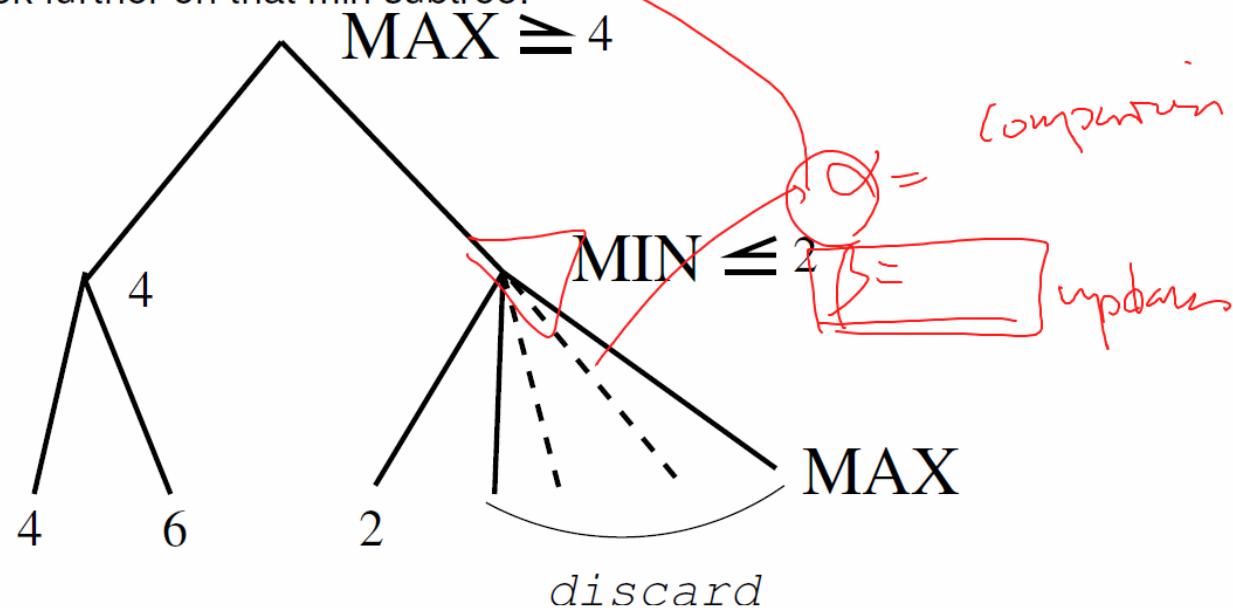


Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.



α Cuts: $\alpha = \text{best MAX so far}$

When the current max value is greater than the successor's min value,
don't look further on that min subtree:

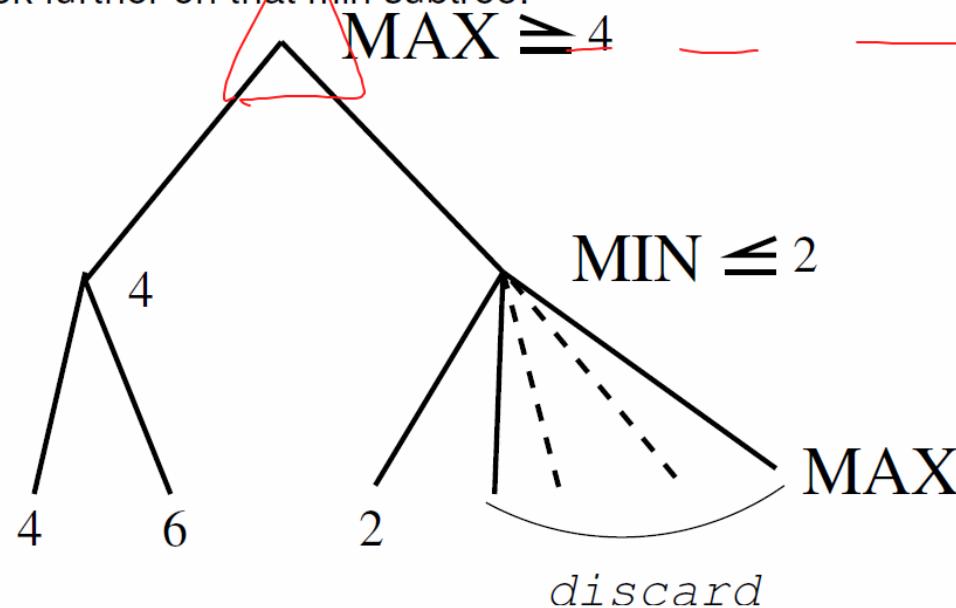


Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.

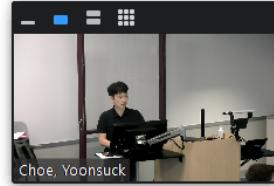


α Cuts: $\alpha = \text{best MAX so far}$

When the current max value is greater than the successor's min value,
don't look further on that min subtree:

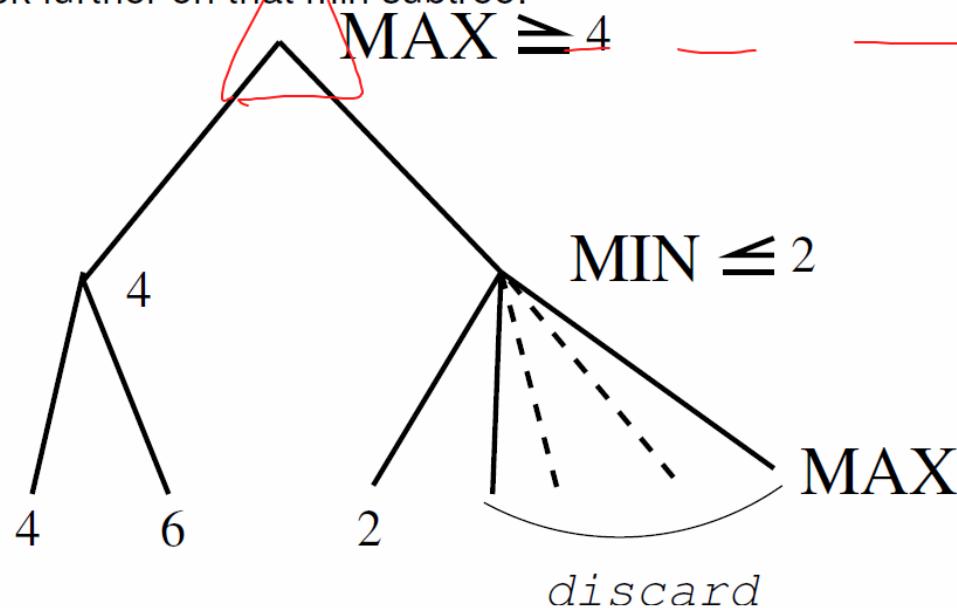


Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.



α Cuts: $\alpha = \text{best MAX so far}$

When the current max value is greater than the successor's min value,
don't look further on that min subtree:

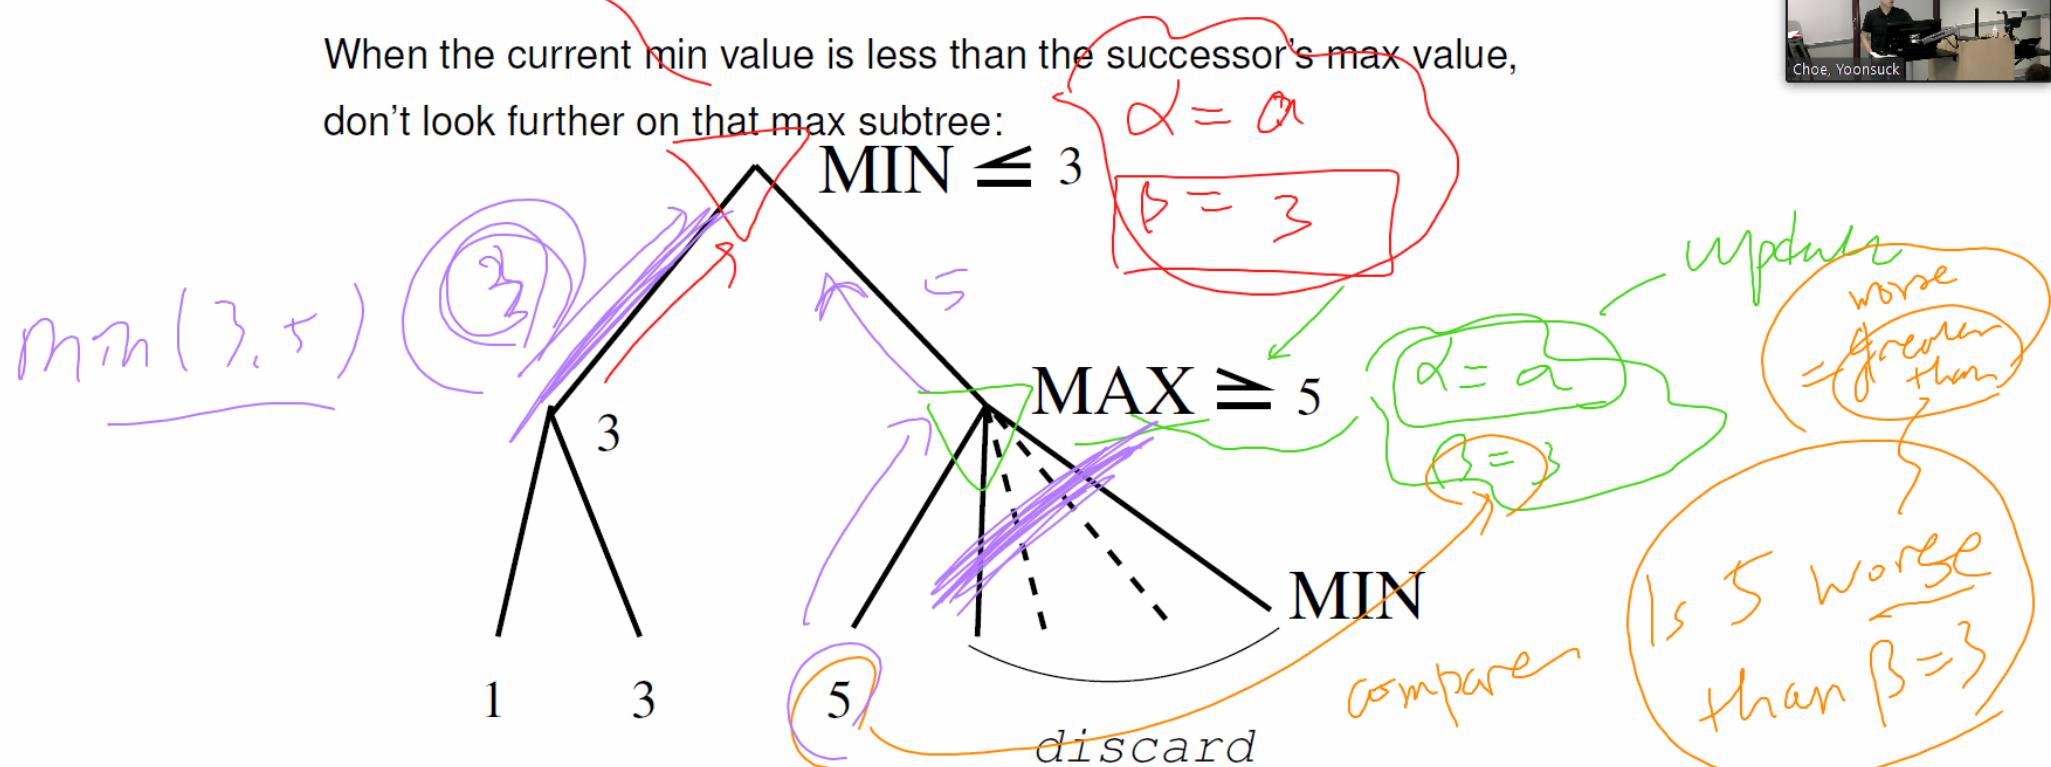


Right subtree can be **at most** 2, so MAX will always choose the left path regardless of what appears next.



β Cuts: $\beta = \text{best MIN so far}$

When the current min value is less than the successor's max value,
don't look further on that max subtree:



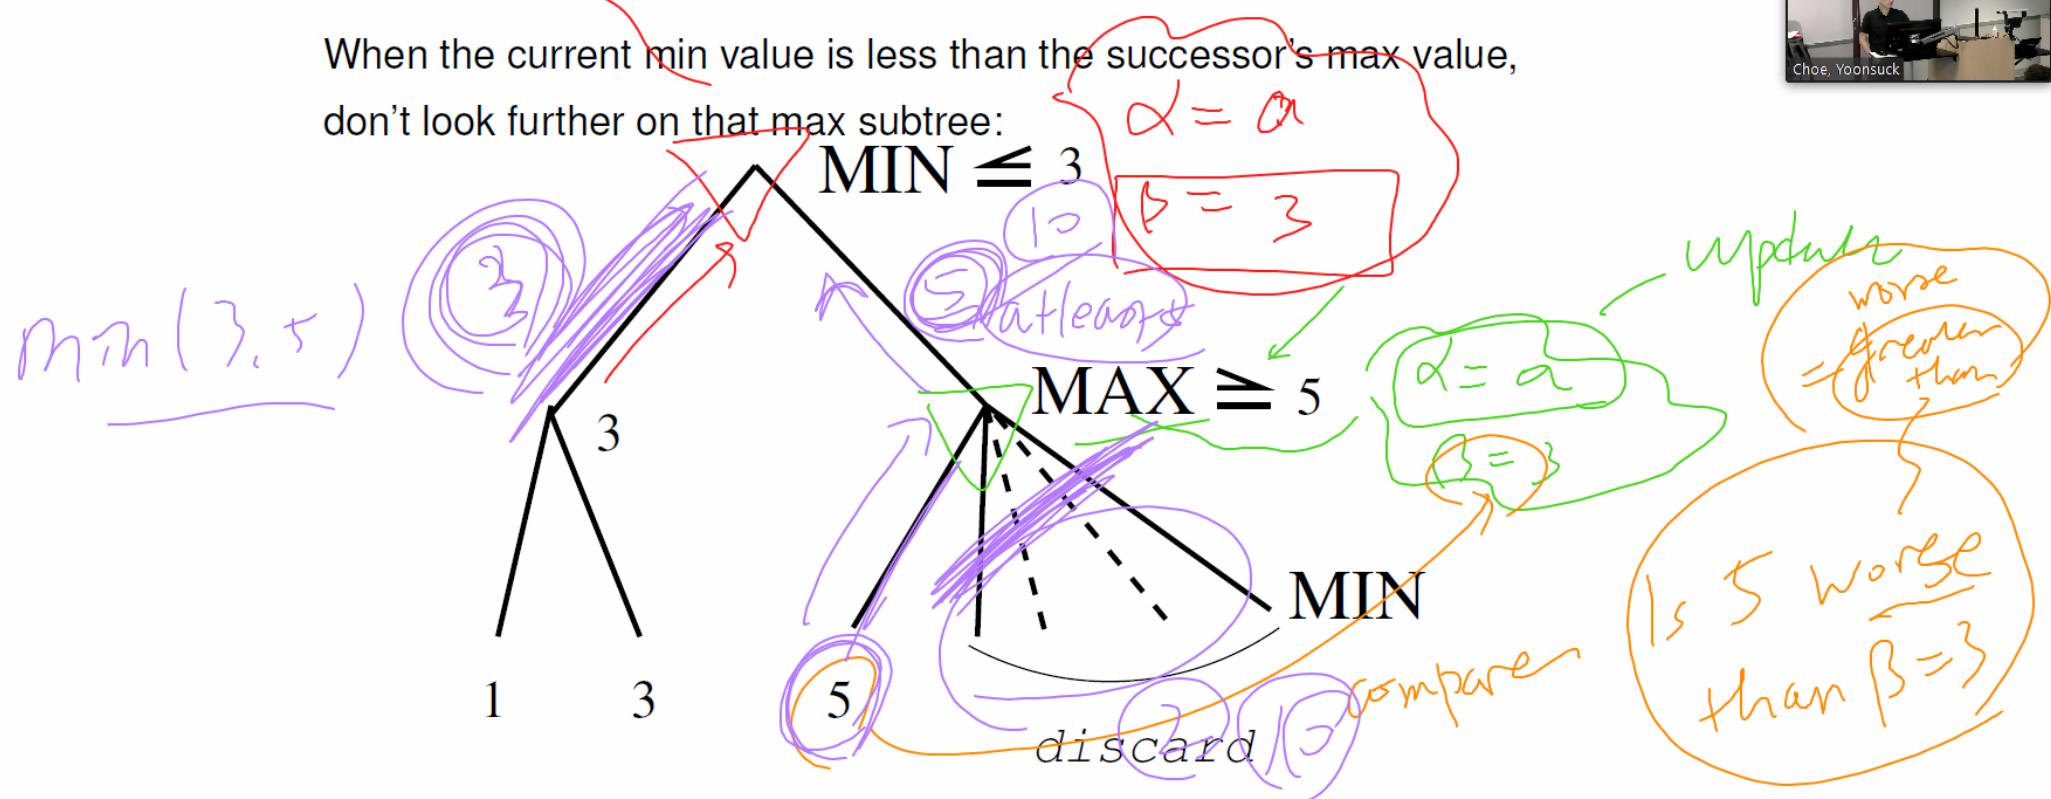
Right subtree can be **at least** 5, so MIN will always choose the left path regardless of what appears next.

β_{\min}



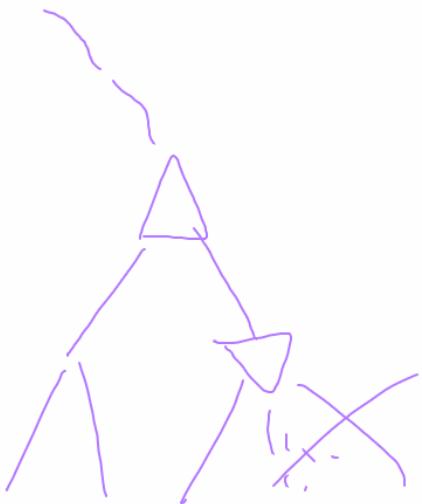
β Cuts: $\beta = \text{best MIN so far}$

When the current min value is less than the successor's max value,
don't look further on that max subtree:

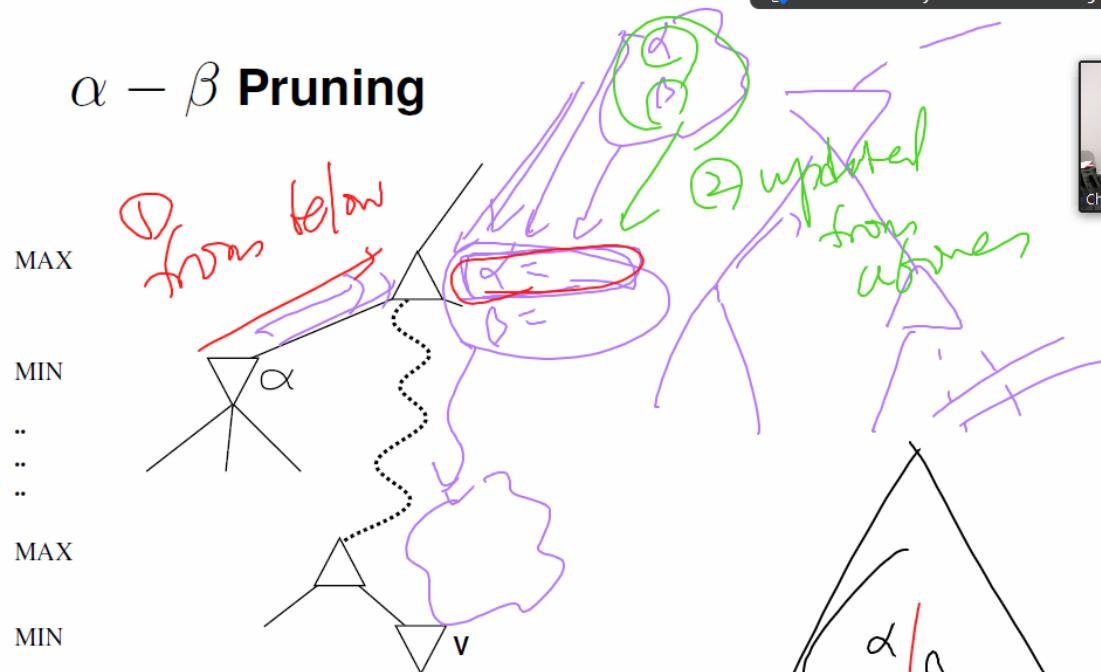


Right subtree can be **at least** 5, so MIN will always choose the left path regardless of what appears next.

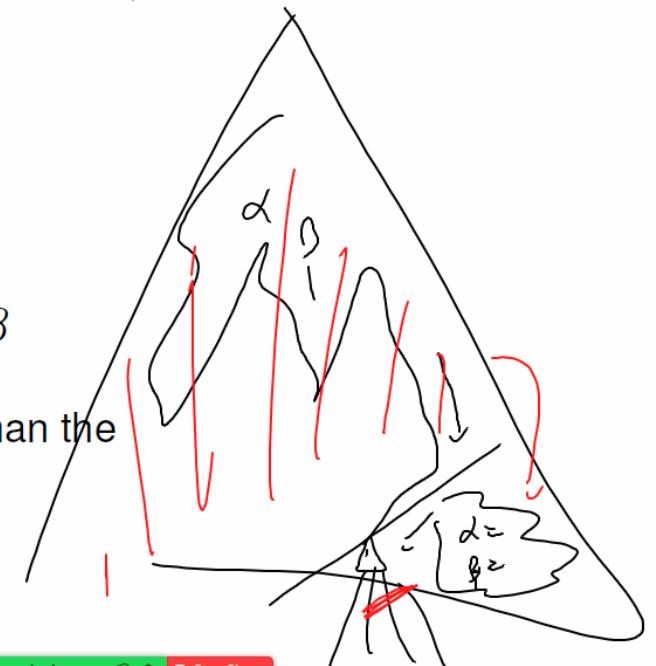
By this point. But time that you encountered this so minimum, 5 value will be returned, and then you mean no component will always just see, I think the other br



$\alpha - \beta$ Pruning

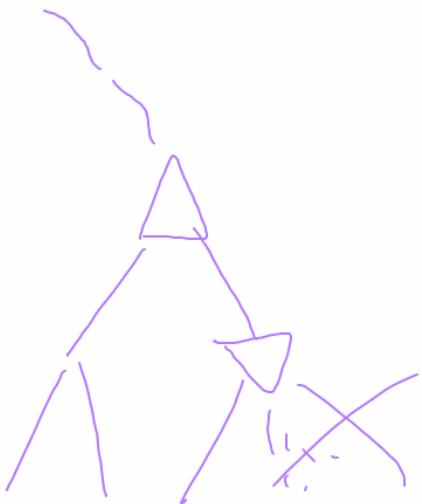


- memory of best MAX value α and best MIN value β
- do not go further on any one that does **worse (!!)** than the remembered α and β

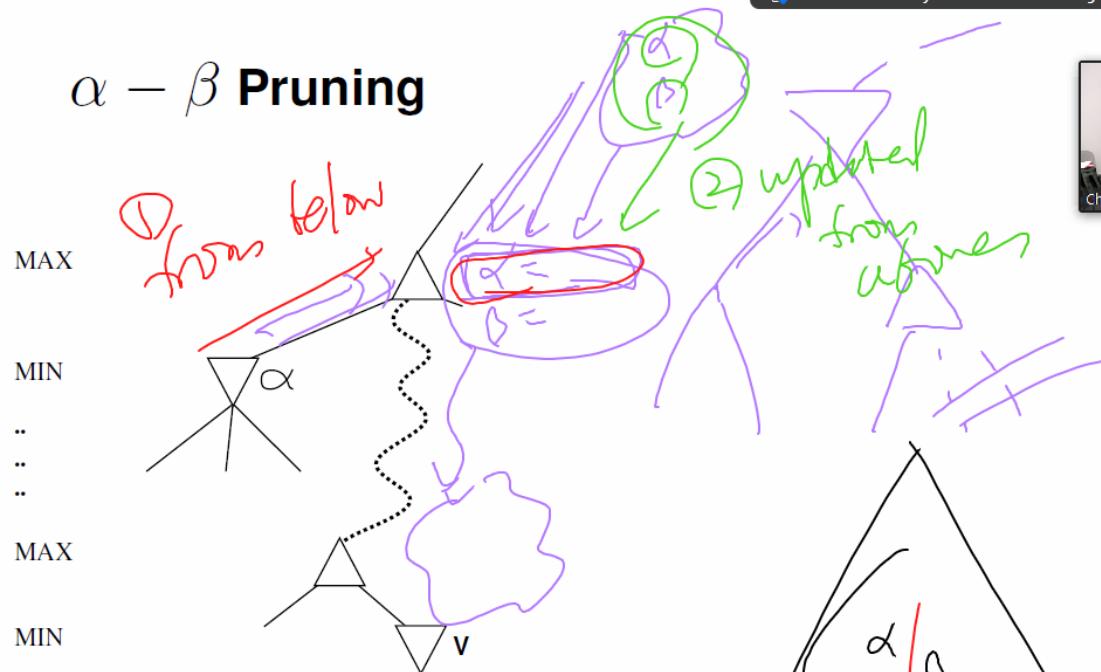


So far in this entire search, and based on that, you make a decision

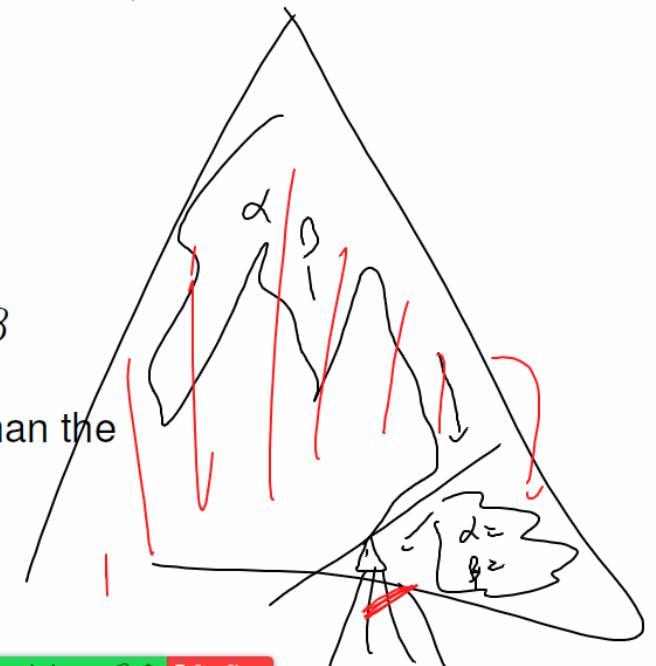




$\alpha - \beta$ Pruning

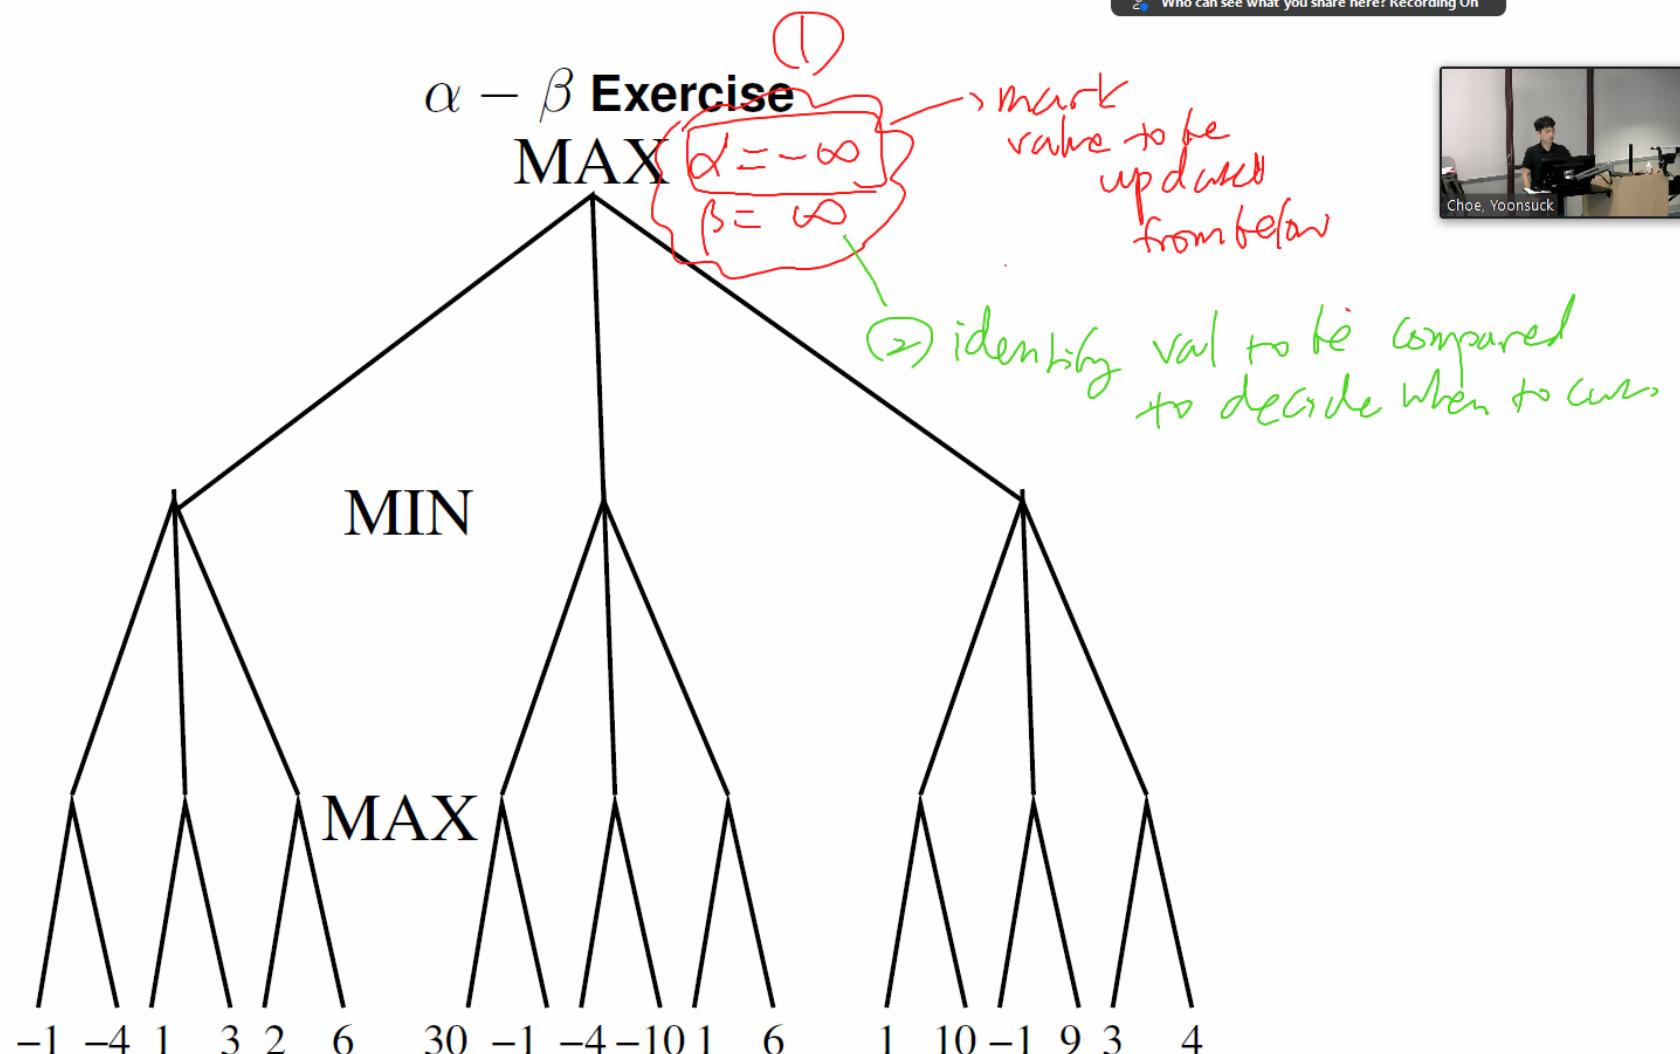


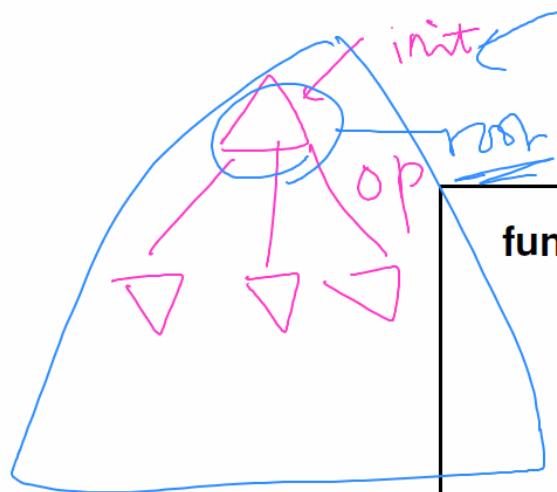
- memory of best MAX value α and best MIN value β
- do not go further on any one that does **worse (!!)** than the remembered α and β



So far in this entire search, and based on that, you make a decision to cut







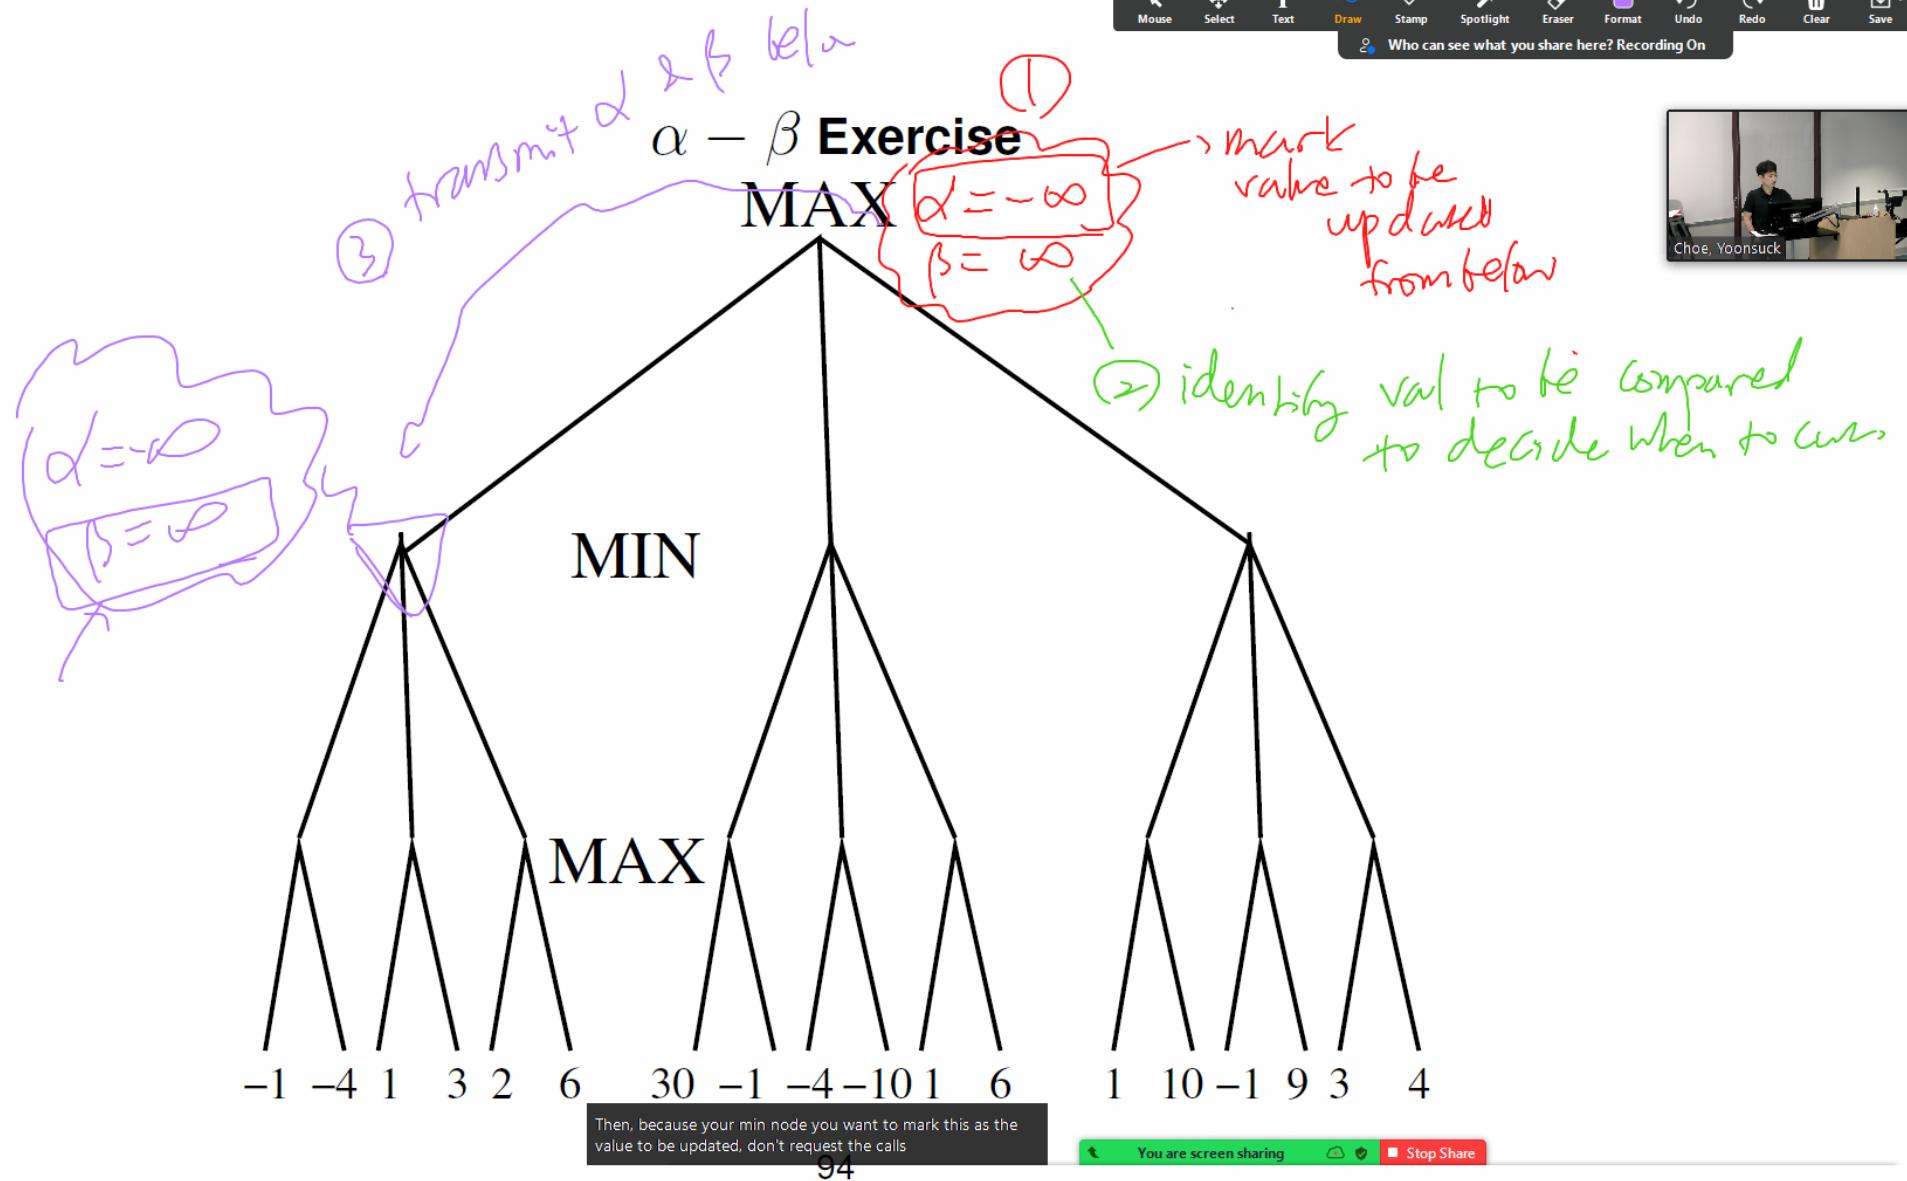
Minimax Decision

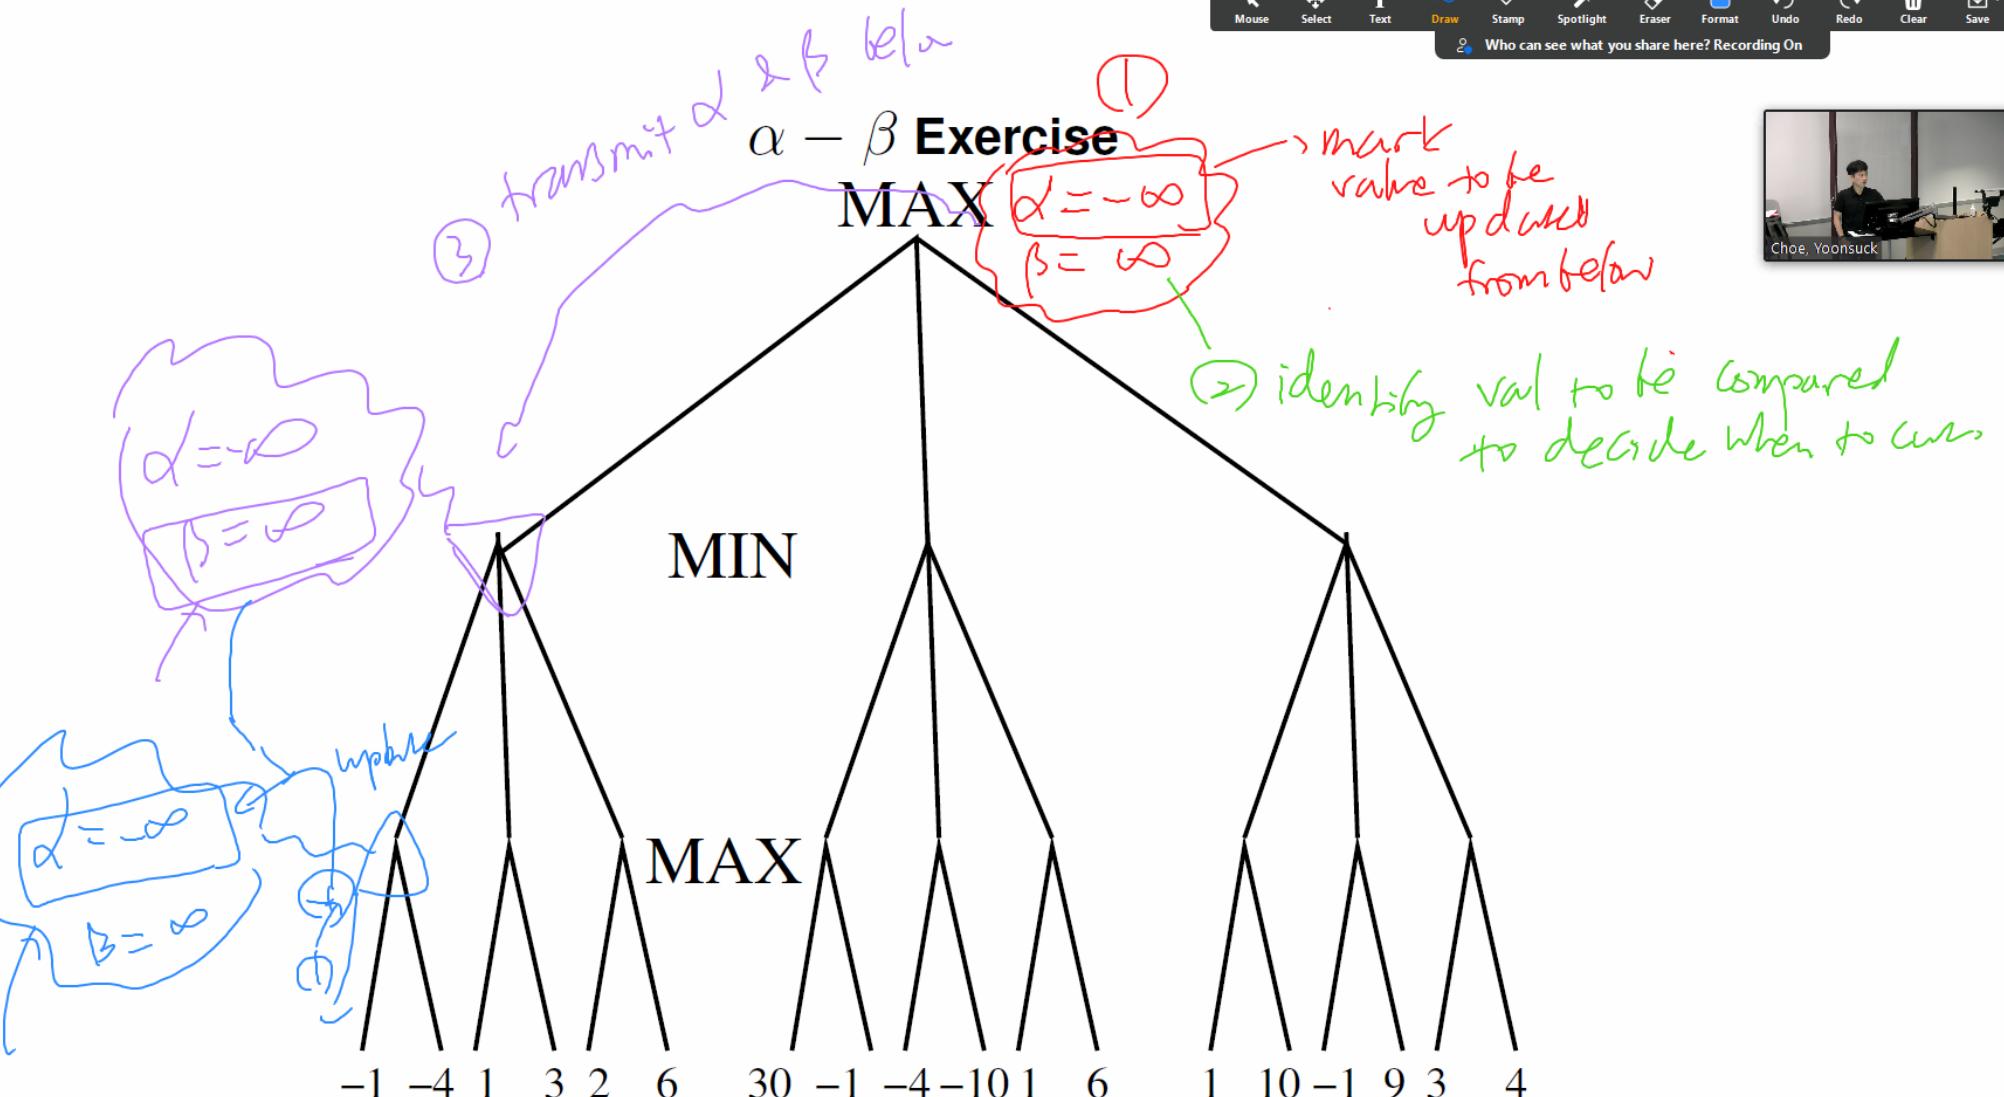
```
function Minimax-Decision(game) returns operator
    return operator that leads to a child state with the
    max(Minimax-Value(child state,game))
```

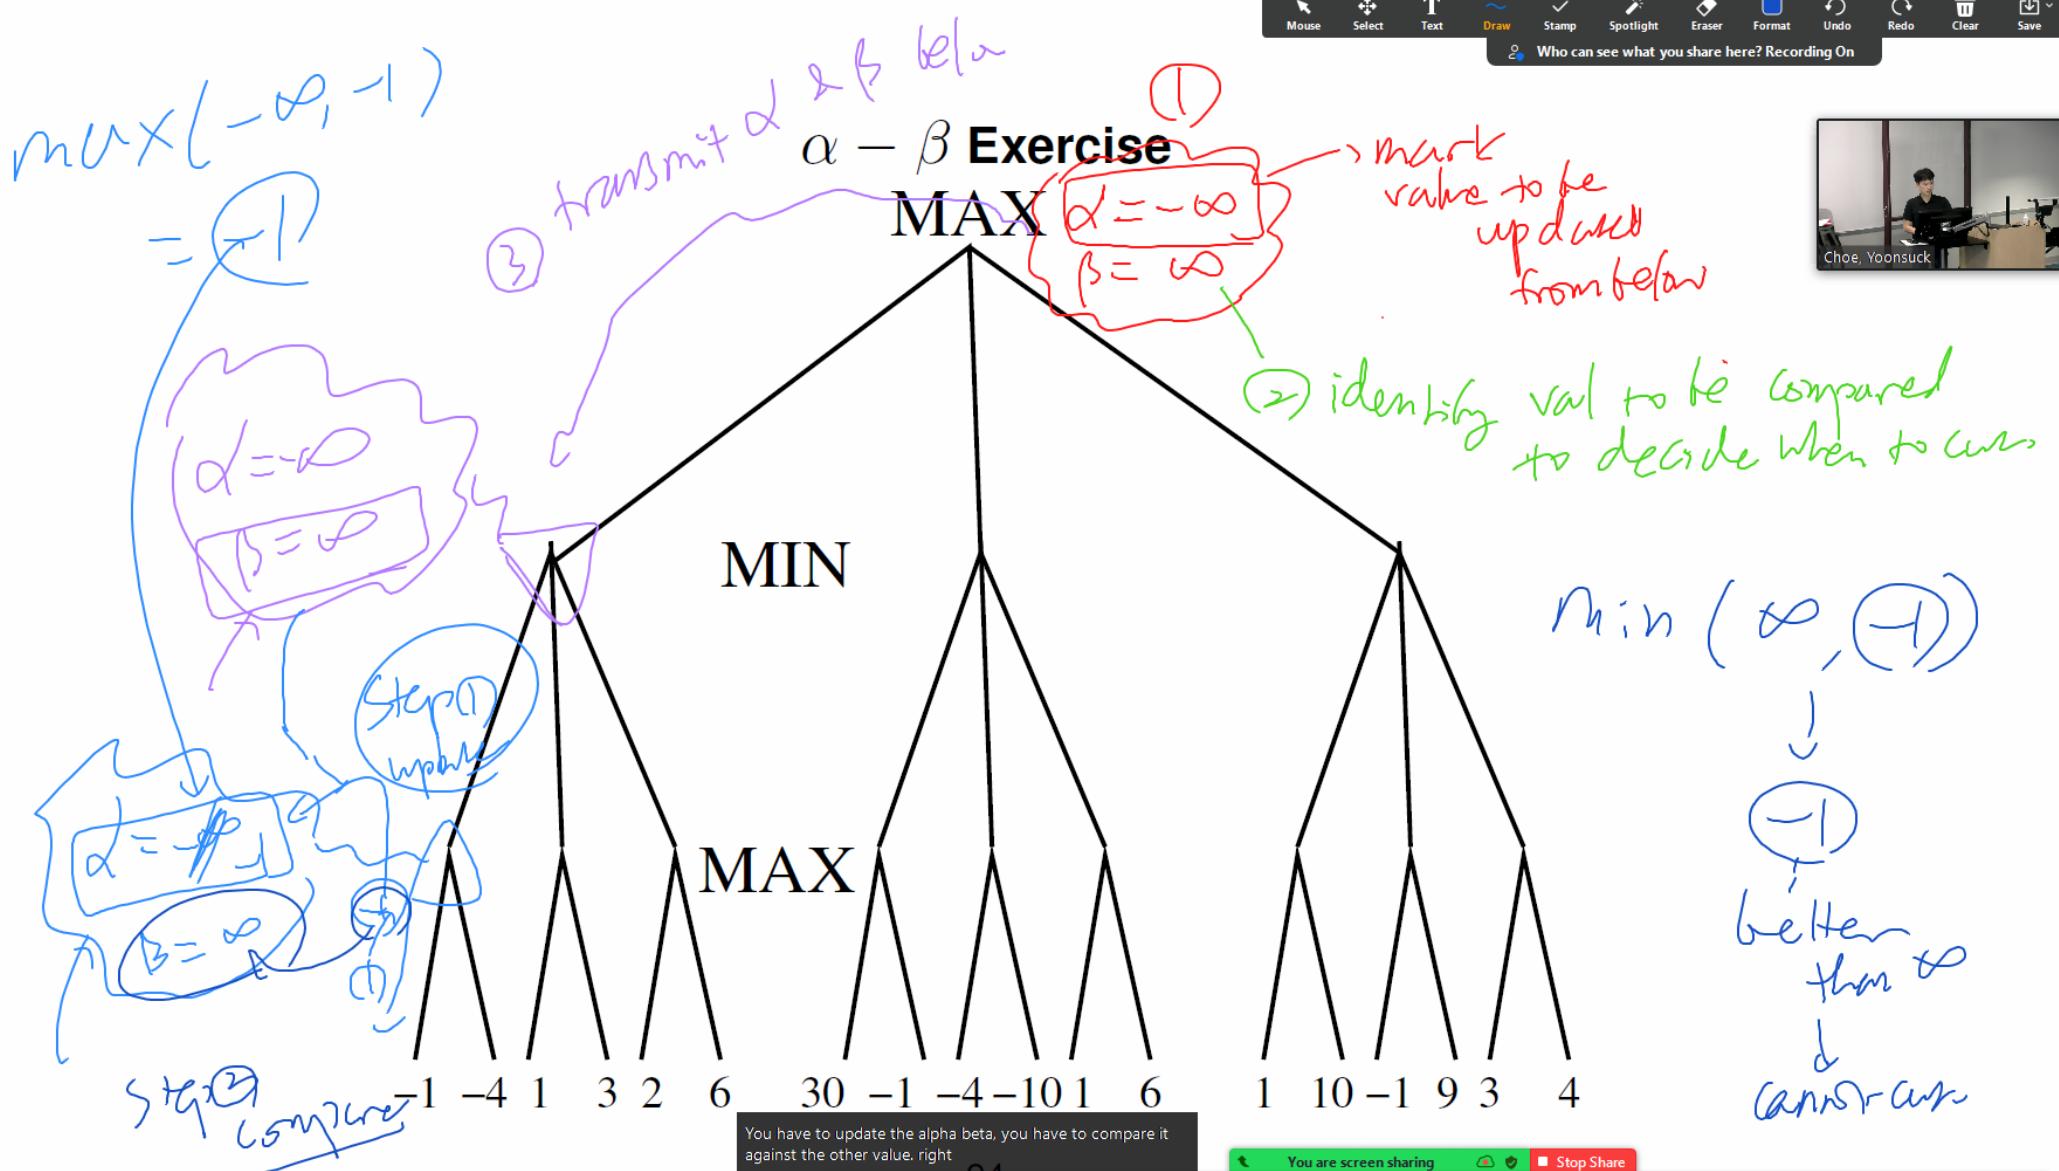
```
function Minimax-Value(state,game) returns utility value
    if Goal(state), return Utility(state)
    else if Max's move then
        → return max of successors' Minimax-Value
    else
        → return min of successors' Minimax-Value
```



Choe, Yoonsuck





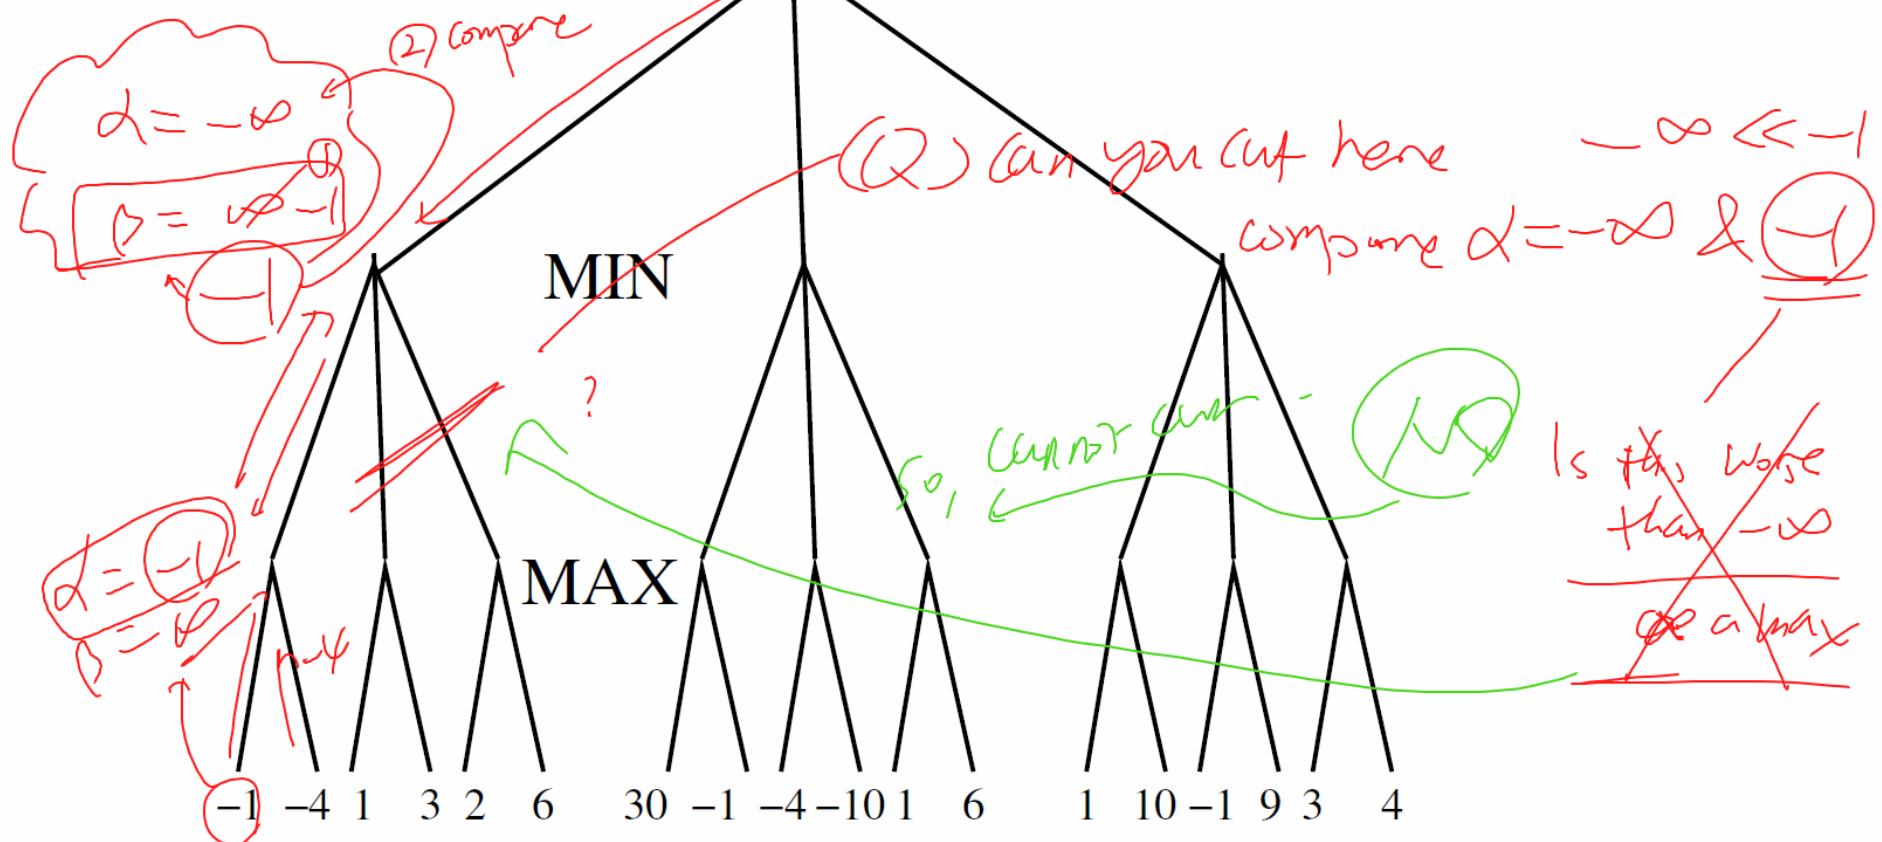




$\alpha - \beta$ Exercise

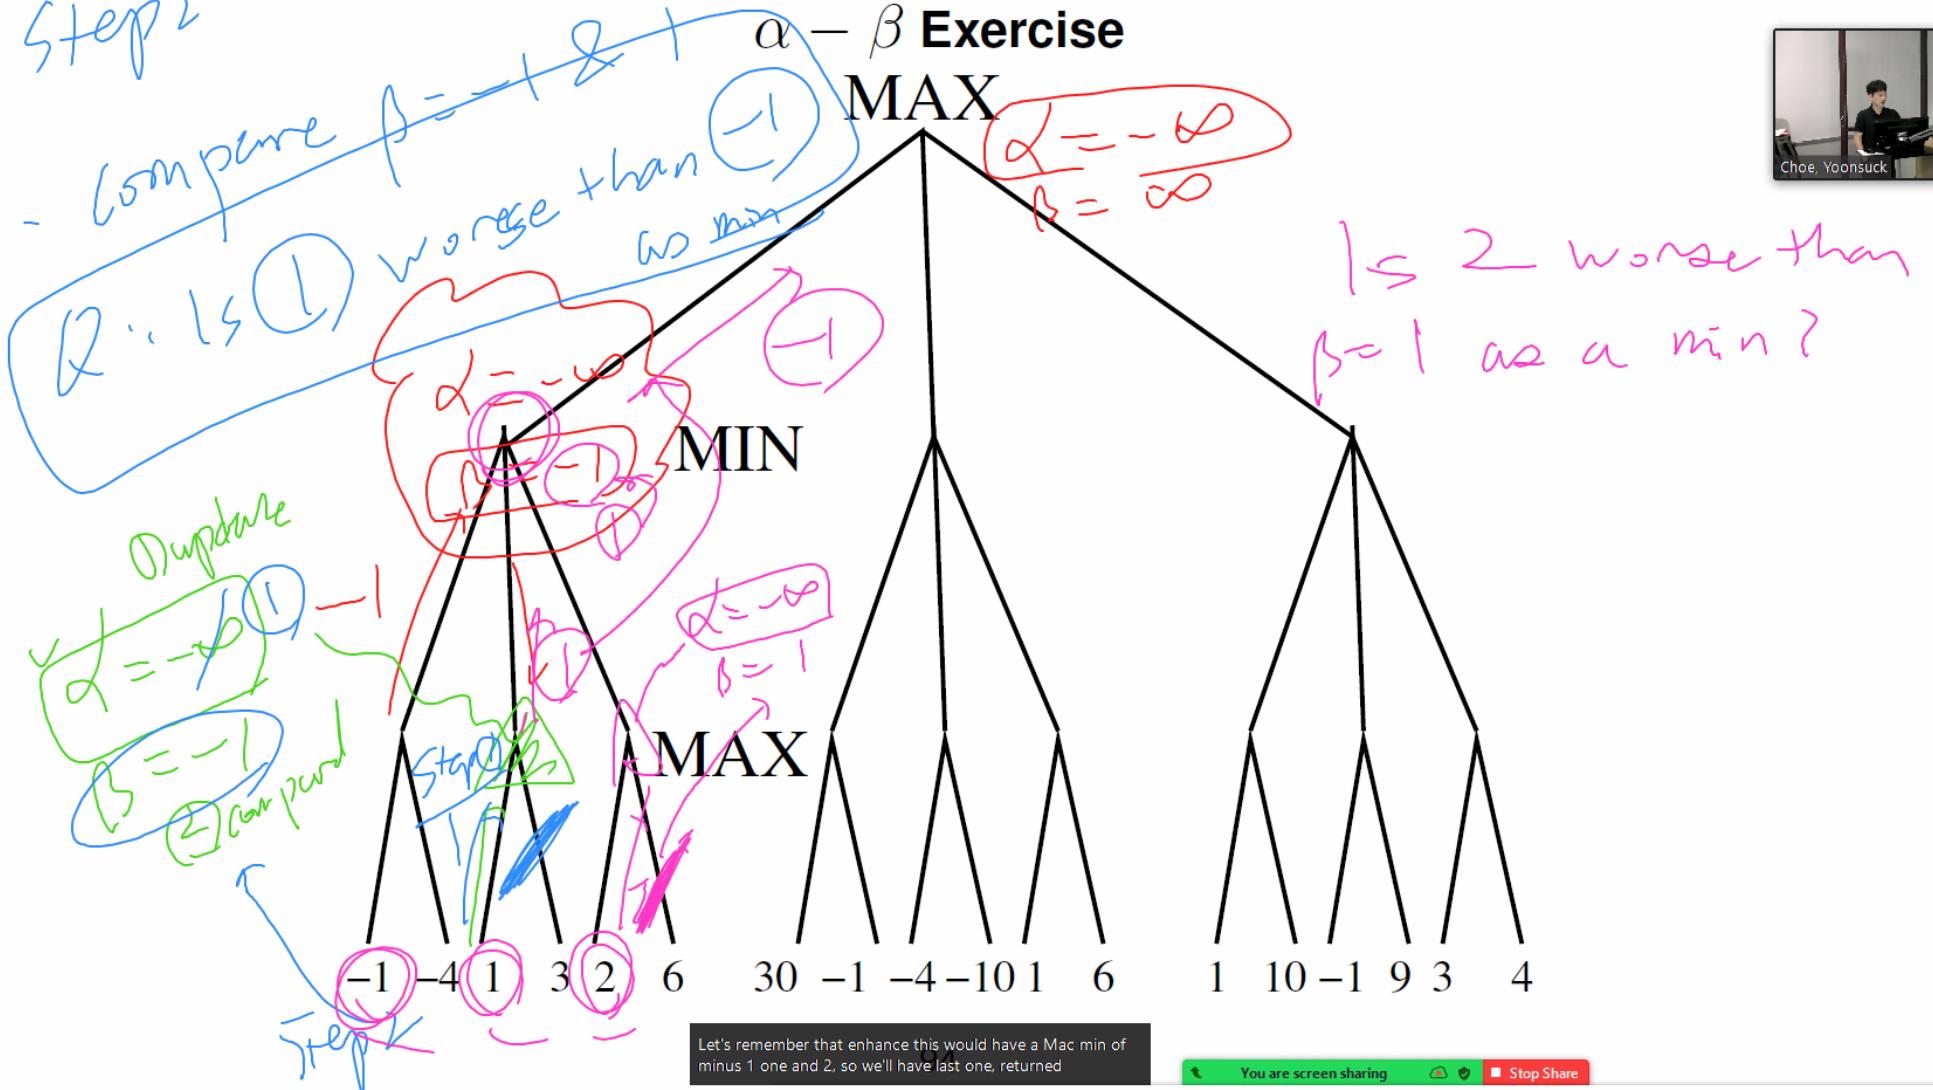
MAX $\alpha = -\infty$

$\beta = \infty$



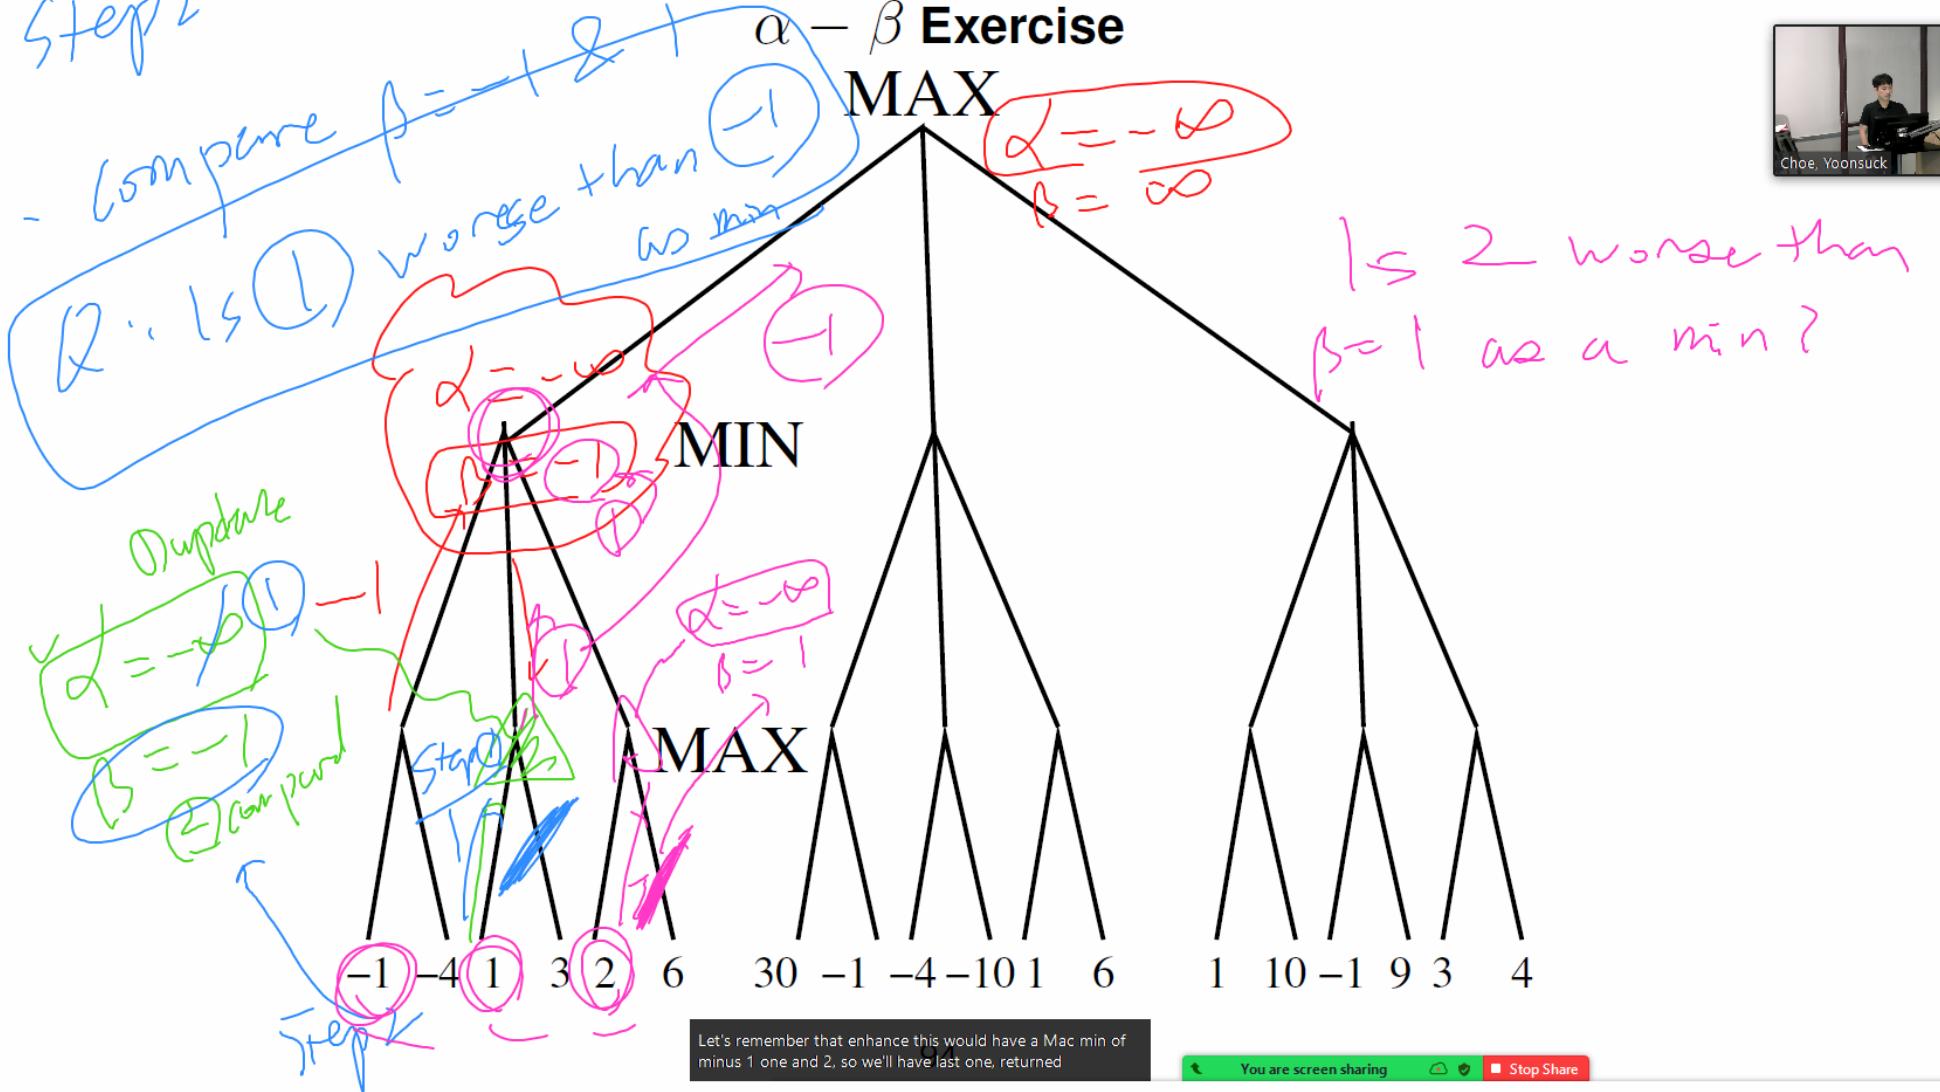


Step 2

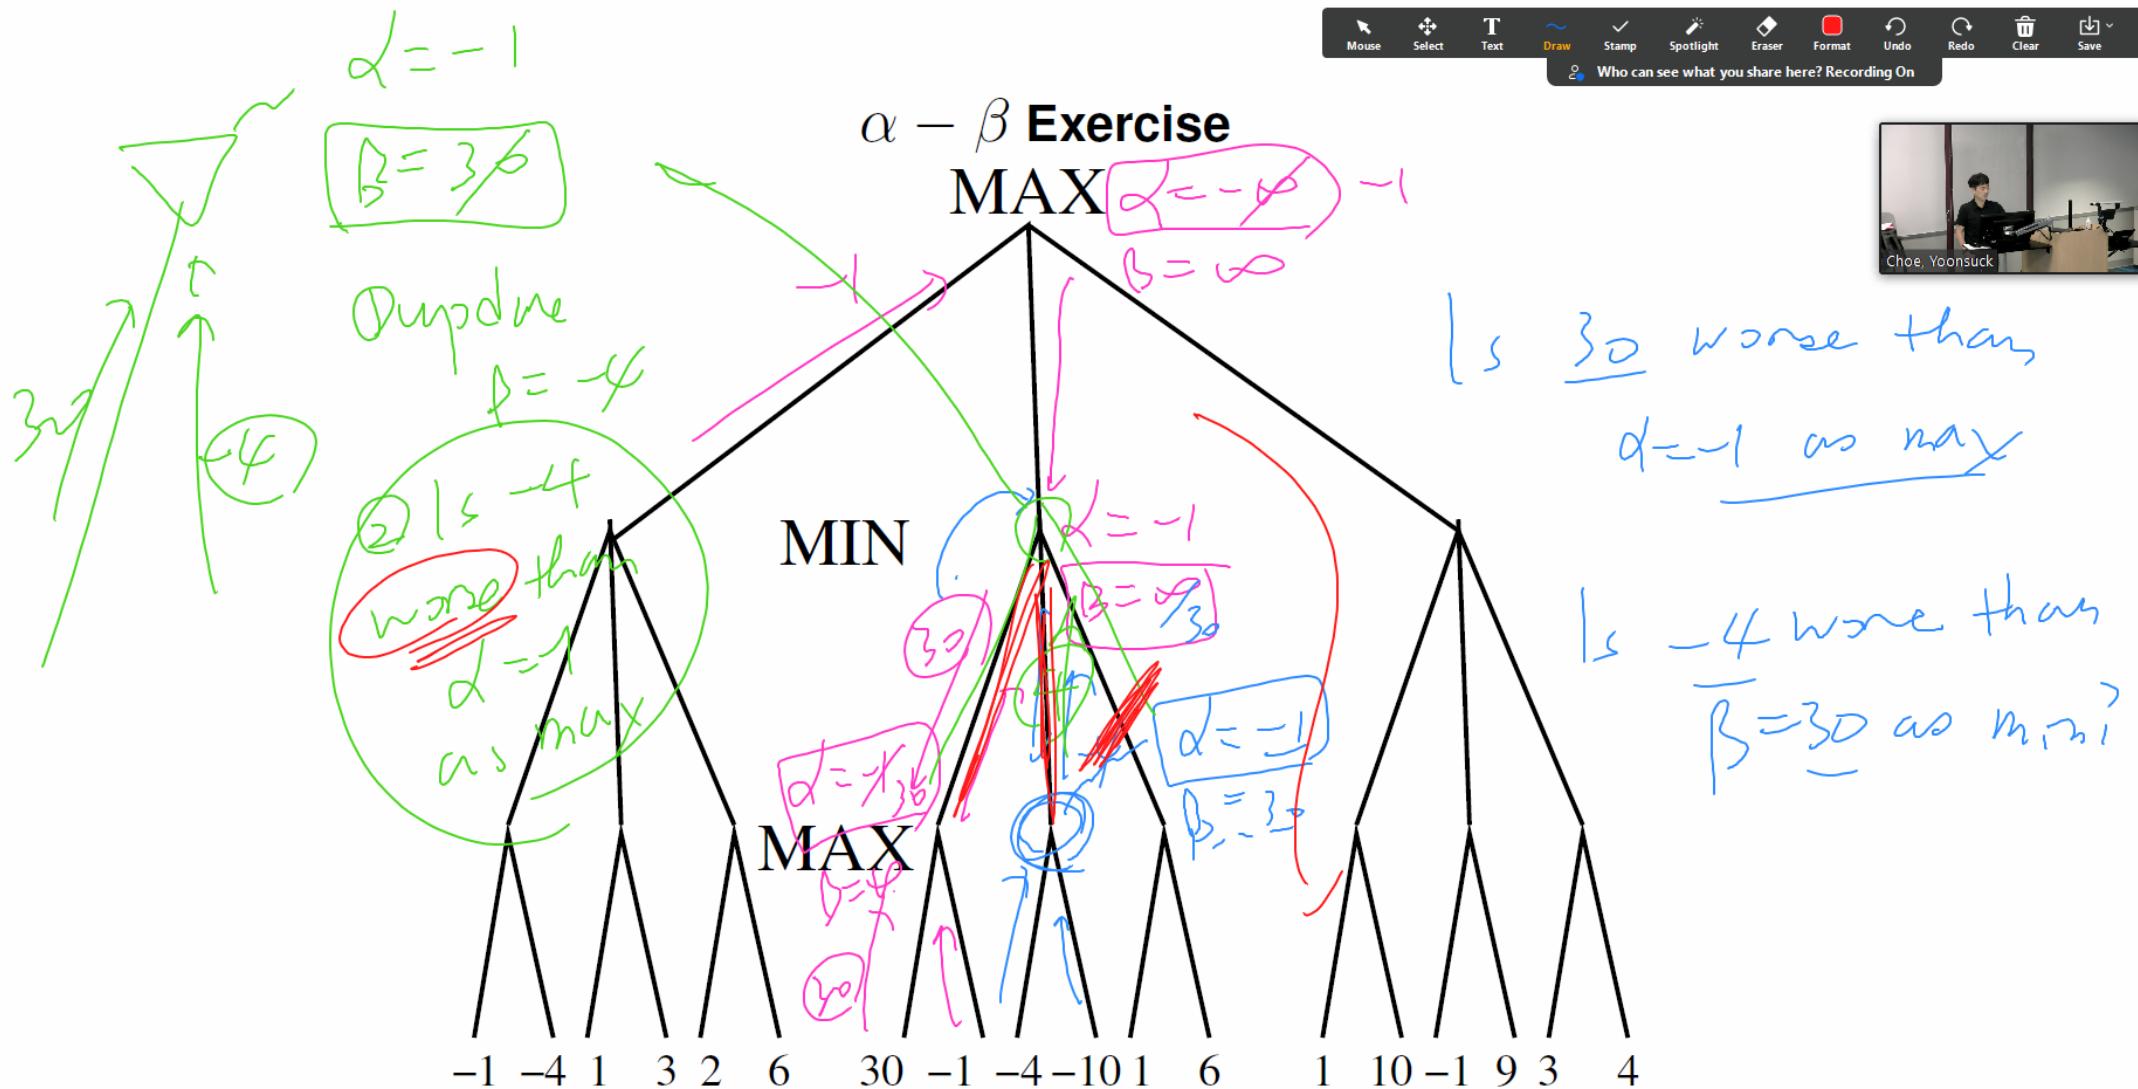




Step 2

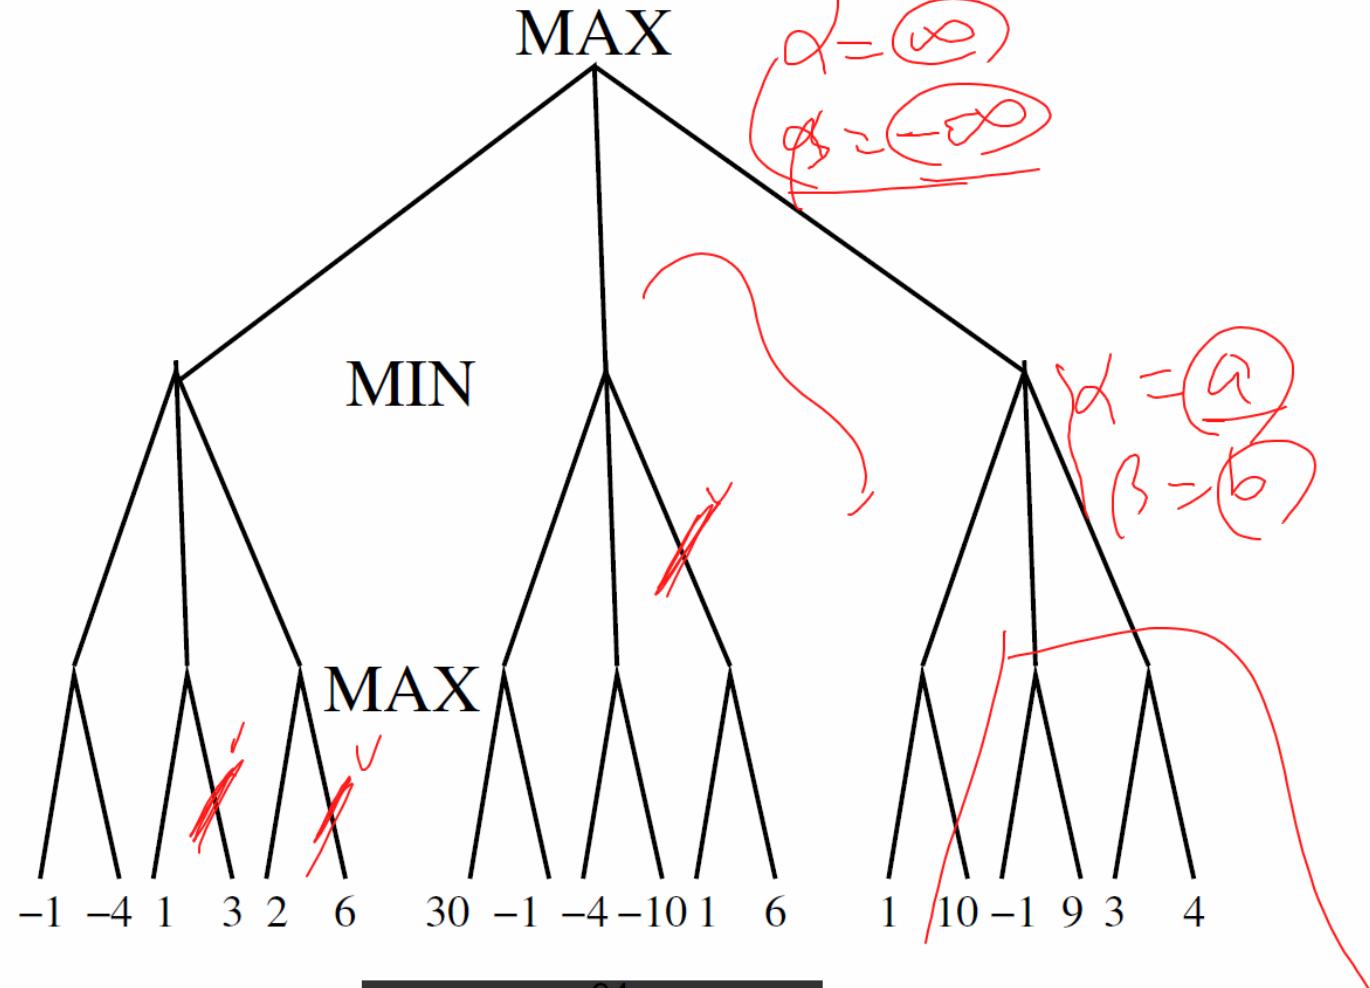


Who can see what you share here? Recording On

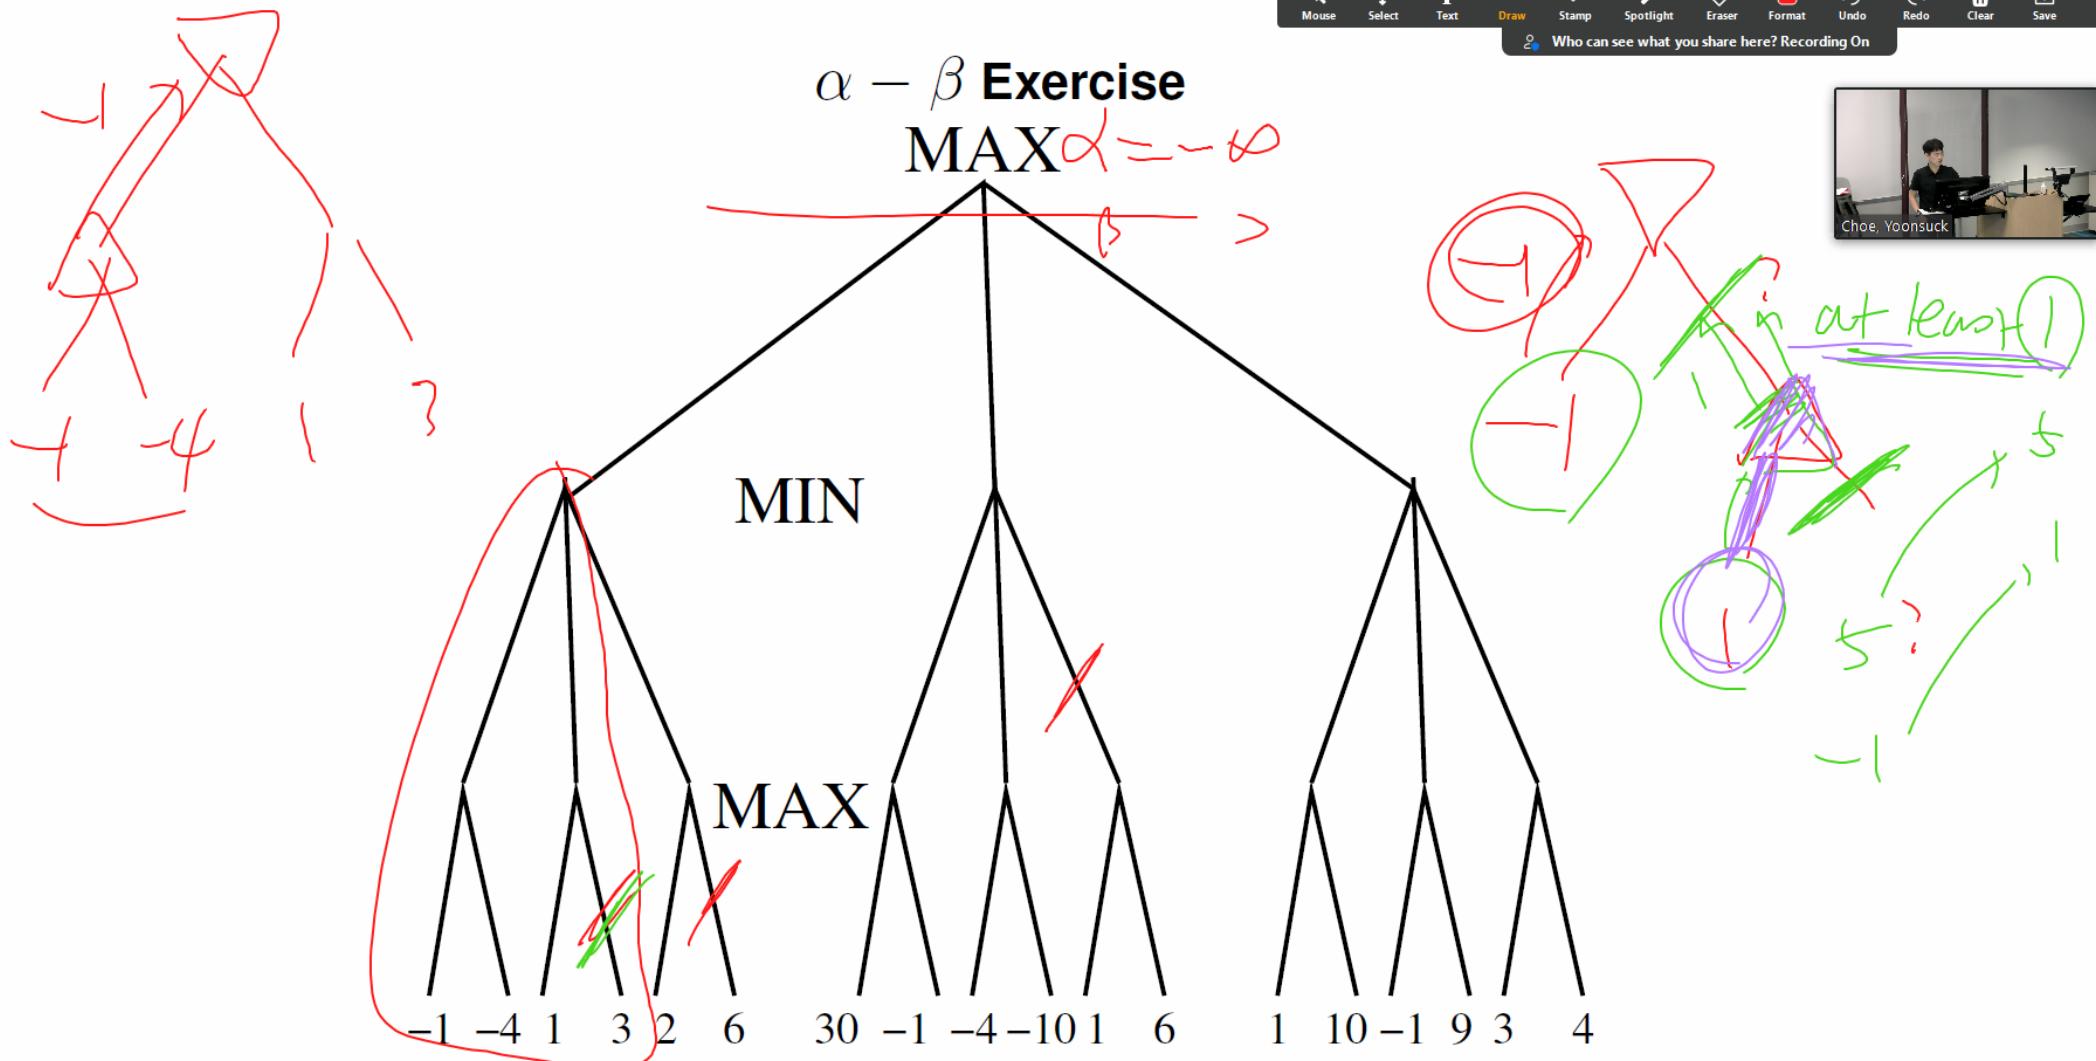




$\alpha - \beta$ Exercise



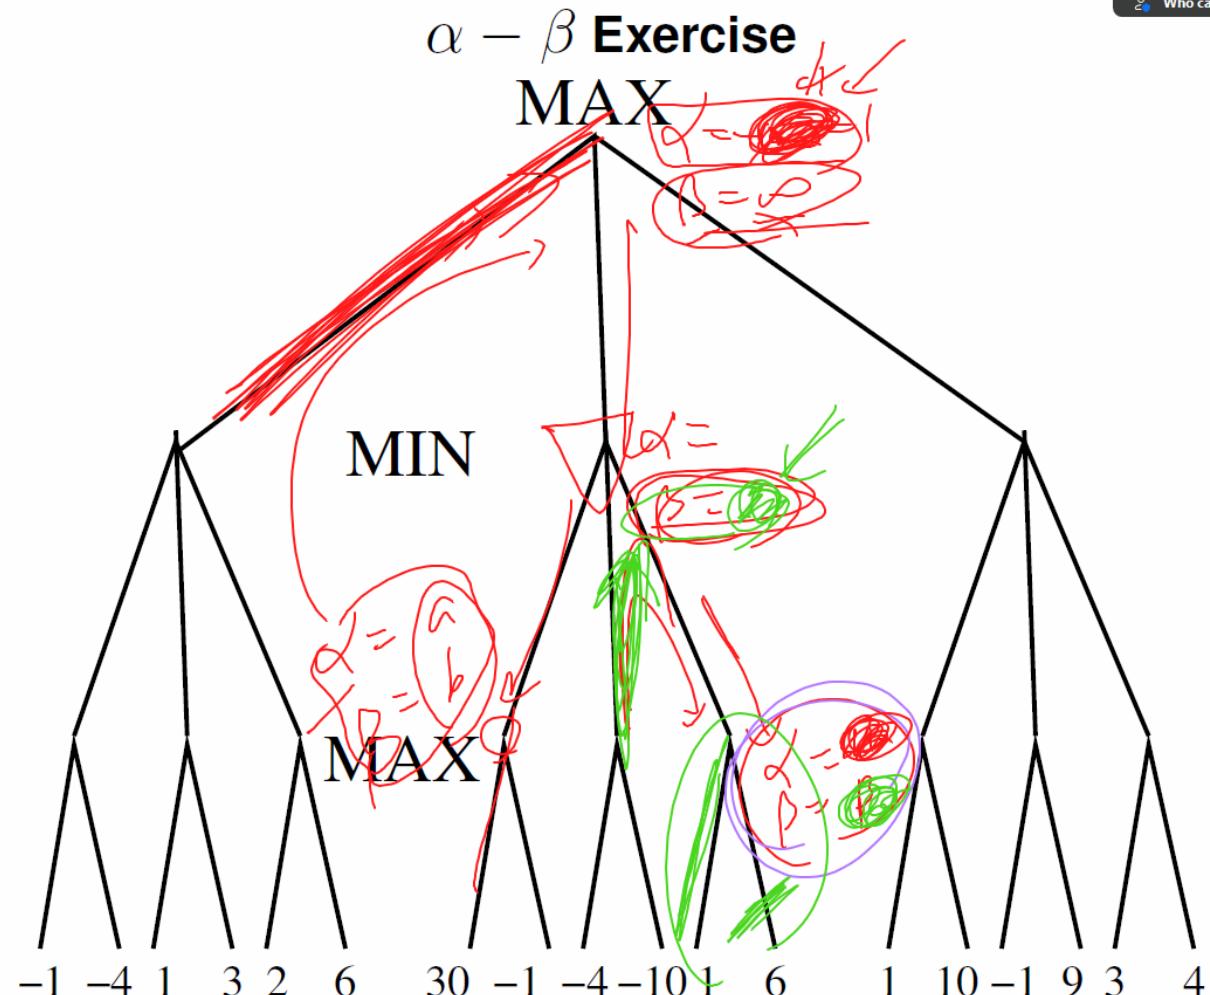
Who can see what you share here? Recording On



So now this is that So you have. The alpha equals minus infinity.

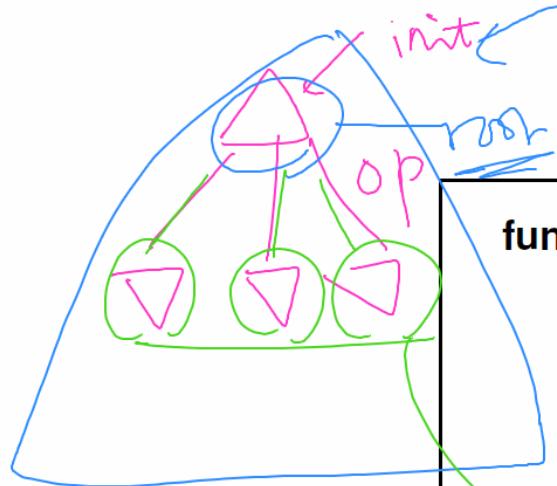
You are screen sharing Stop Share

Who can see what you share here? Recording On



you can start cutting alright, but you have to actually look
at that to make to make any decision about that

You are screen sharing



Minimax Decision

```
function Minimax-Decision(game) returns operator
    return operator that leads to a child state with the
    max(Minimax-Value(child state,game))
```

```
function Minimax-Value(state,game) returns utility value
    if Goal(state), return Utility(state)
    else if Max's move then
        → return max of successors' Minimax-Value
    else
        → return min of successors' Minimax-Value
```

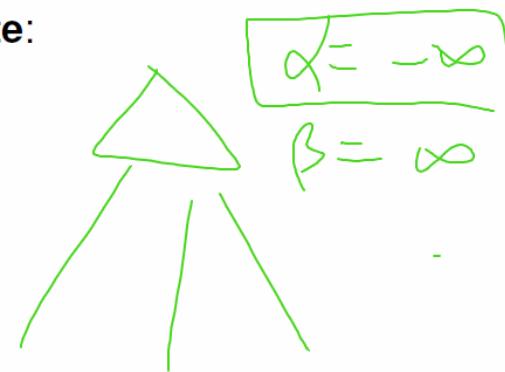


$\alpha - \beta$ Pruning: Initialization



Along the path from the beginning to the current **state**:

- α : best MAX value
 - initialize to $-\infty$
- β : best MIN value
 - initialize to ∞



$\alpha - \beta$ Pruning Tips

- At a MAX node:

- Only α is updated with the MAX of successors.
- Cut is done by checking if returned $v \geq \beta$.
- If all fails, $\text{MAX}(v \text{ of successors})$ is returned.

- At a MIN node:

- Only β is updated with the MIN of successors.
- Cut is done by checking if returned $v \leq \alpha$.
- If all fails, $\text{MIN}(v \text{ of successors})$ is returned.



Choe, Yoonsuck

$\alpha - \beta$ Pruning Tips

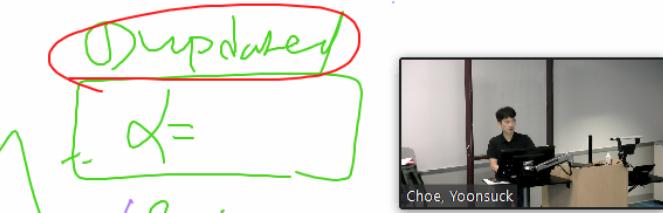
- At a MAX node:

- Only α is updated with the MAX of successors.
- Cut is done by checking if returned $v \geq \beta$.
- If all fails, MAX(v of successors) is returned.

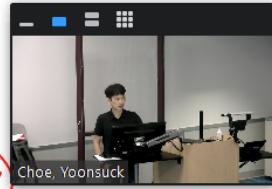
- At a MIN node:

- Only β is updated with the MIN of successors.
- Cut is done by checking if returned $v \leq \alpha$.
- If all fails, MIN(v of successors) is returned.

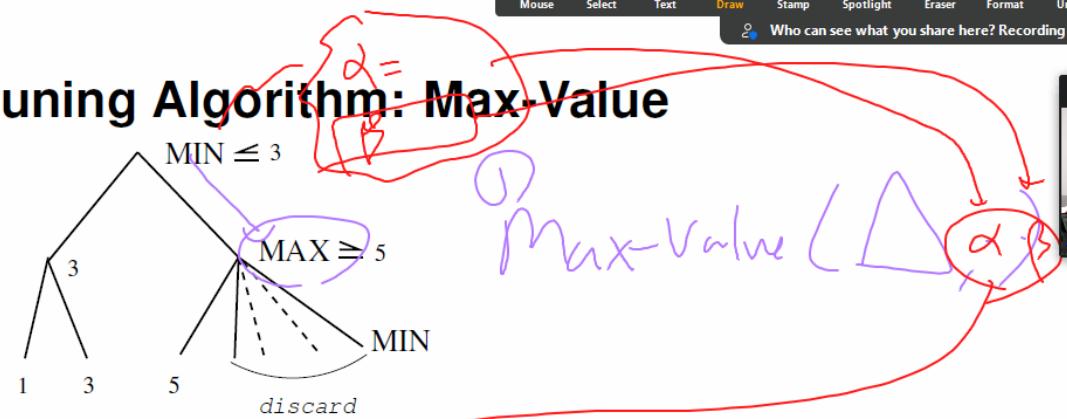
Min Alpha is the best, Max. So if you just keep this in mind, then understanding the algorithm is a participant



(1) Updated
 $\alpha =$
 β
 (2) Compared
 for a cut.
 β_{cur}
 is v worse than β
 (as a min)
 ~ don't
 is v worse than
 (α_{max})



$\alpha - \beta$ Pruning Algorithm: Max-Value



```

function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility(state)
 $v \leftarrow -\infty$ 
for each s in Successor(state) do
     $v \leftarrow \text{Max}(v, \text{Min-Value}(s, \text{game}, \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$  /*  $\beta$  CUT!! */
     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
end
return  $v$ 

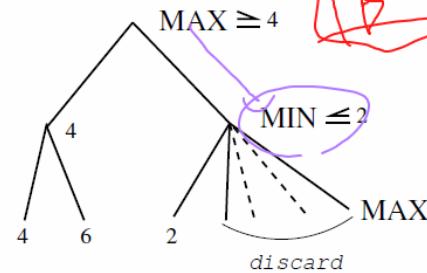
```

Then this often beta value gets transmitted to you through this requested call

You are screen sharing Stop Share



$\alpha - \beta$ Pruning Algorithm: Min-Value



```

function Min-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility (state)
     $v \leftarrow \infty$ 
    for each s in Successor(state) do
        •  $v \leftarrow \text{Min}(v, \text{Max-Value}(s, \text{game}, \alpha, \beta))$ 
        • if  $v \leq \alpha$  then return  $v$       /*  $\alpha$  CUT!! */
        •  $\beta \leftarrow \text{Min}(\beta, v)$ 
    end
    return  $v$ 

```



$\alpha - \beta$ Pruning Algorithm: Min-Value



```

function Min-Value (state, game,  $\alpha, \beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility (state)
 $v \leftarrow \infty$ 
for each s in Successor(state) do
     $v \leftarrow \text{Min}(v, \text{Max-Value}(s, game, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$  /*  $\alpha$  CUT!! */
     $\beta \leftarrow \text{Min}(\beta, v)$ 
end
return  $v$ 

```

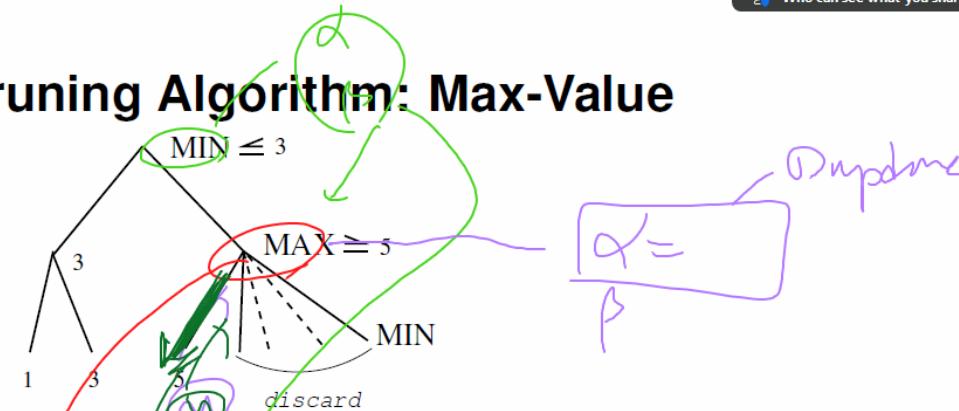
Sorry on a minnode, and then you still have the same alpha
and beta value transmitted from your current or your
opponent through this

100

You are screen sharing Stop Share



$\alpha - \beta$ Pruning Algorithm: Max-Value



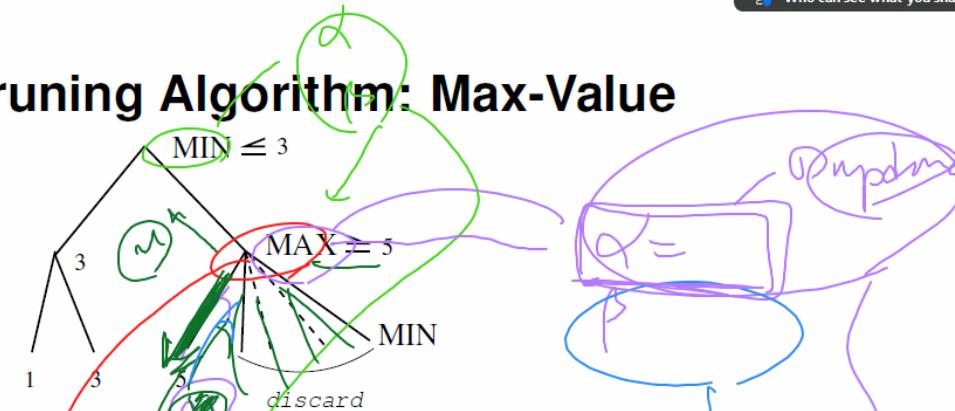
```

function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility(state)
 $v \leftarrow -\infty$ 
for each  $s$  in Successor(state) do recursive call
    •  $v \leftarrow \text{Max}(v, \text{Min-Value}(s, game, \alpha, \beta))$ 
    • if  $v \geq \beta$  then return  $v$  /*  $\beta$  CUT!! */
    •  $\alpha \leftarrow \text{Max}(\alpha, v)$ 
end
return  $v$ 

```



$\alpha - \beta$ Pruning Algorithm: Max-Value



```

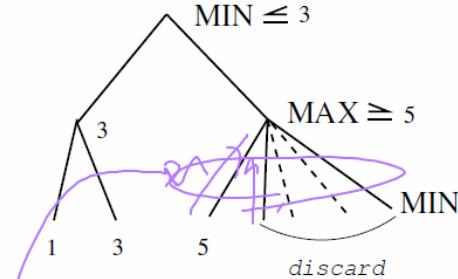
function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility(state)
     $v \leftarrow -\infty$ 
    for each s in Successor(state) do recursive call
         $v \leftarrow \text{Max}(v, \text{Min-Value}(s, game, \alpha, \beta))$ 
        if  $v \geq \beta$  then return v /*  $\beta$  CUT!! */
         $\alpha \leftarrow \text{Max}(\alpha, v)$ 
    end
    return v

```

So far is rules then cut. Okay, and also the update would happen for the particular value of a beta value that you marked putting to the particular type of your note



$\alpha - \beta$ Pruning Algorithm: Max-Value



```

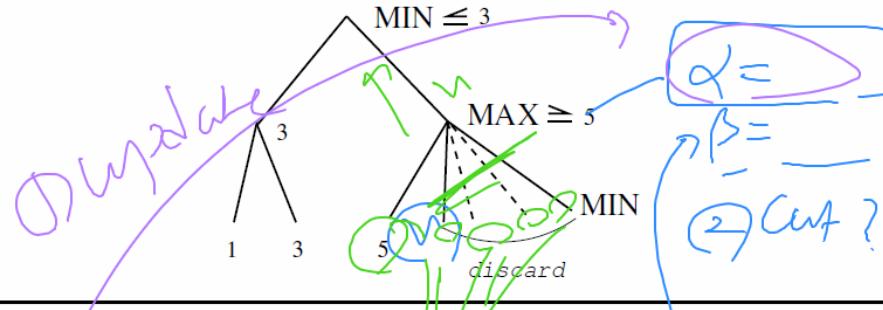
function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
     $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
    if Cutoff(state) then return Utility(state)
     $v \leftarrow -\infty$ 
    for each  $s$  in Successor(state) do
         $v \leftarrow \text{Max}(v, \text{Min-Value}(s, game, \alpha, \beta))$ 
        if  $v \geq \beta$  then return  $v$  /*  $\beta$  CUT!! */
         $\alpha \leftarrow \text{Max}(\alpha, v)$ 
    end
    return  $v$ 

```

So far as you stole the branches, so these mean value functions will be called on all



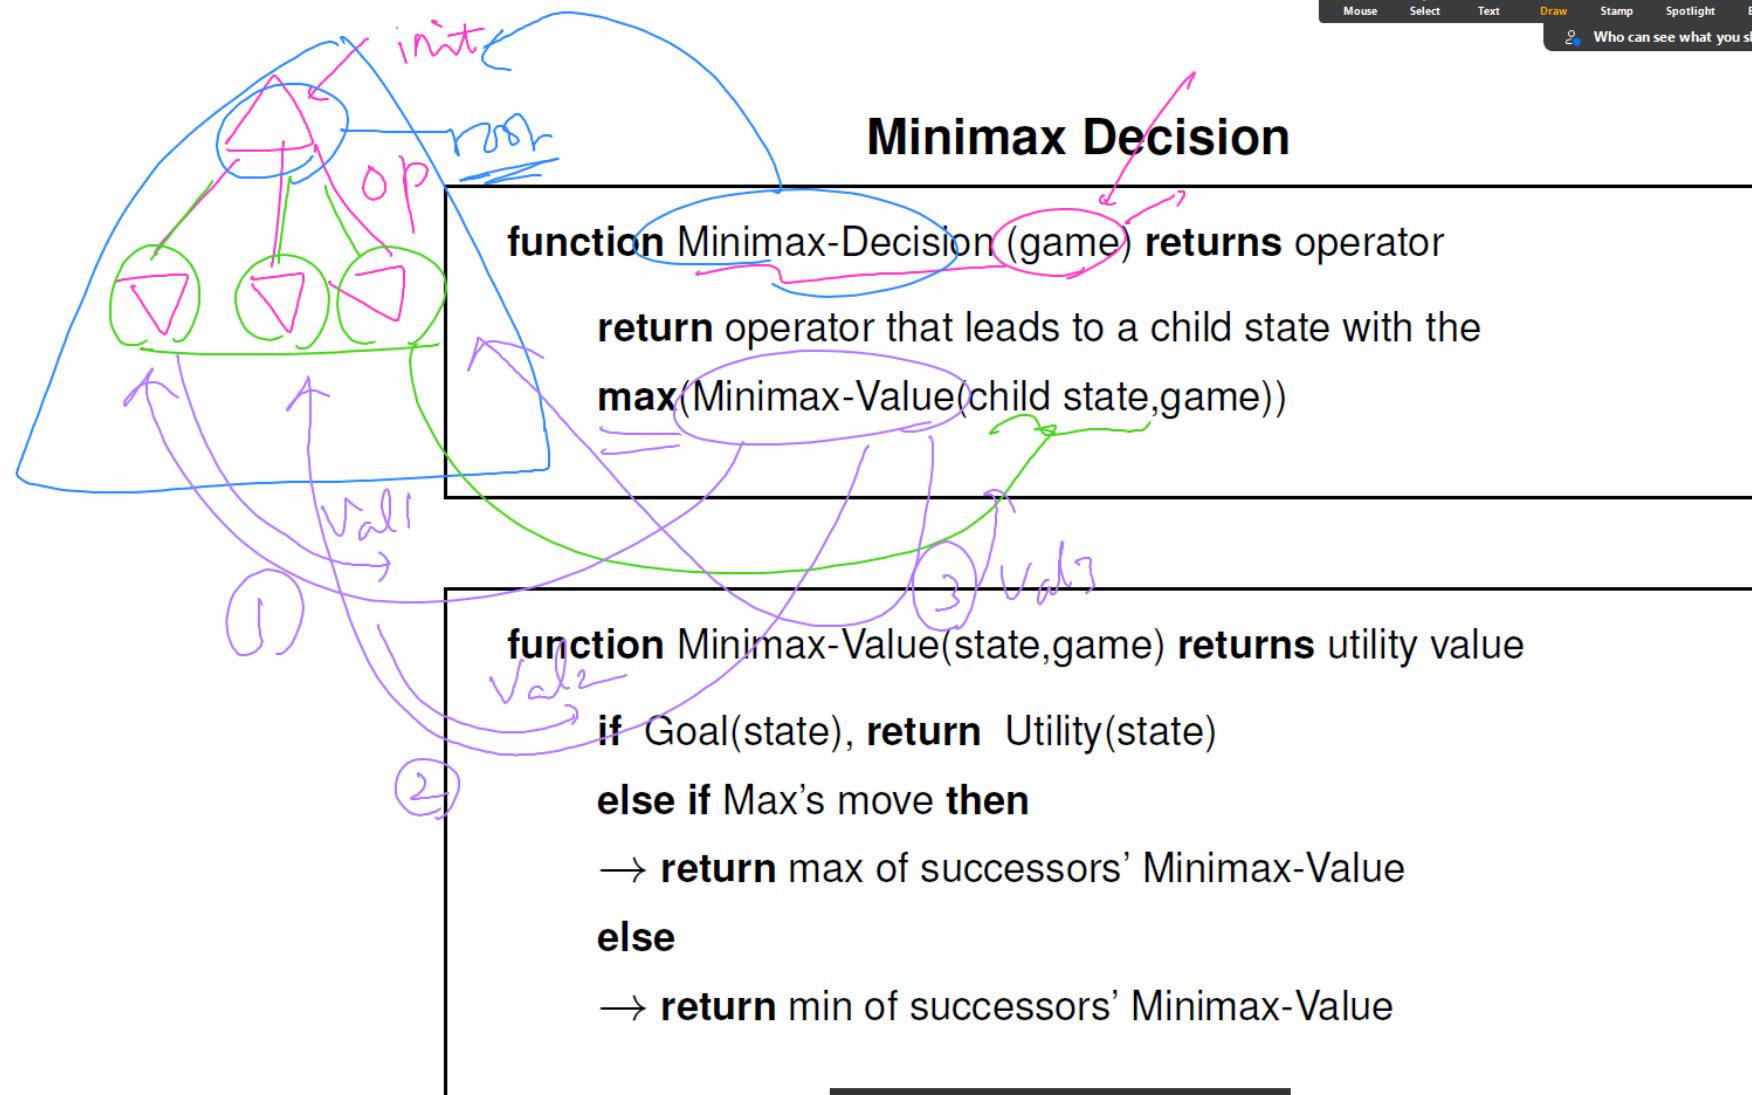
$\alpha - \beta$ Pruning Algorithm: Max-Value



```

function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
     $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
    if Cutoff(state) then return Utility(state)
     $v \leftarrow -\infty$ 
    for each s in Successor(state) do
         $v \leftarrow \text{Max}(v, \text{Min-Value}(s, game, \alpha, \beta))$ 
        if  $v \geq \beta$  then return v /*  $\beta$  CUT!! */
         $\alpha \leftarrow \text{Max}(\alpha, v)$ 
    end
    return v
  
```

recursive call.



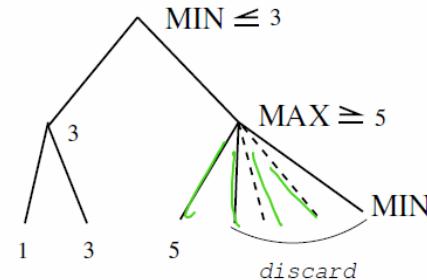
I'm sorry, child, which are written you under the value then now, you just return the Max with the operator that gives gave you the Max

You are screen sharing Stop Share

Max-value

$\alpha - \beta$ Pruning Algorithm: Max-Value

→ min value
→ min value



```

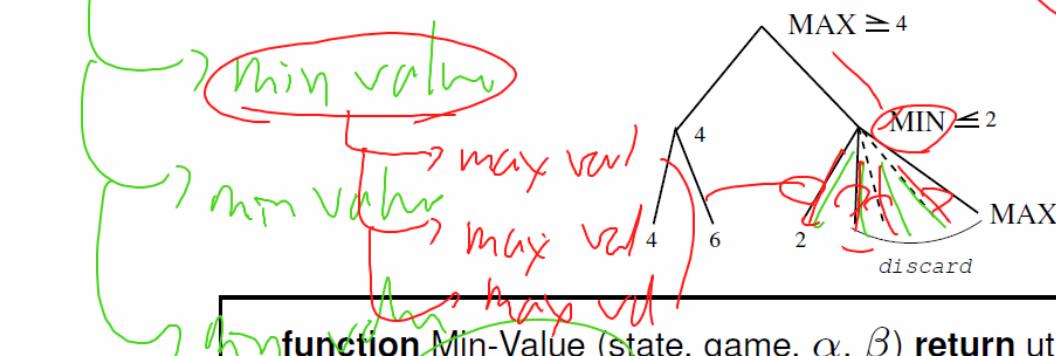
function Max-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
     $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
    if Cutoff(state) then return Utility(state)
     $v \leftarrow -\infty$ 
    for each  $s$  in Successor(state) do
        •  $v \leftarrow \text{Max}(v, \text{Min-Value}(s, game, \alpha, \beta))$ 
        • if  $v \geq \beta$  then return  $v$  /*  $\beta$  CUT!! */
        •  $\alpha \leftarrow \text{Max}(\alpha, v)$ 
    end
    return  $v$ 

```



MAX-value

$\alpha - \beta$ Pruning Algorithm: Min-Value



```

function Min-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
     $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
    if Cutoff(state) then return Utility (state)
     $v \leftarrow \infty$ 
    for each  $s$  in Successor(state) do
        •  $v \leftarrow \text{Min}(v, \text{Max-Value}(s, game, \alpha, \beta))$ 
        • if  $v \leq \alpha$  then return  $v$  /*  $\alpha$  CUT!! */
        •  $\beta \leftarrow \text{Min}(\beta, v)$ 
    end
    return  $v$ 

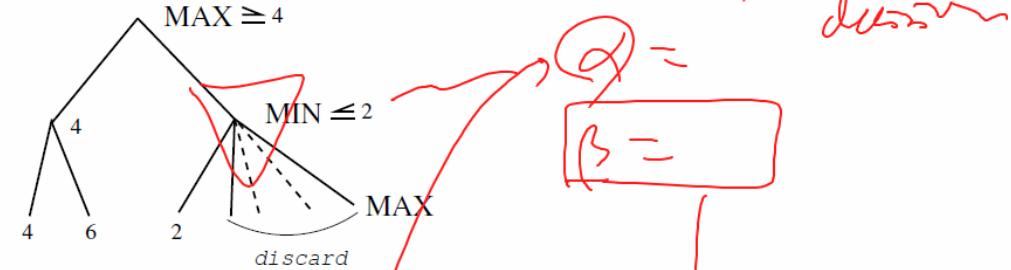
```

So all of these would be these

100

You are screen sharing Stop Share

$\alpha - \beta$ Pruning Algorithm: Min-Value



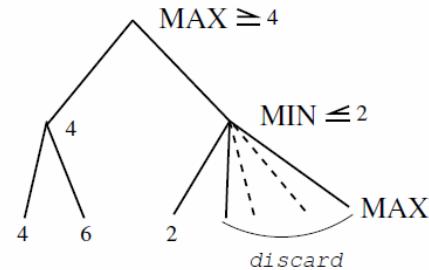
```

function Min-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility (state)
 $v \leftarrow \infty$ 
for each  $s$  in Successor(state) do
    •  $v \leftarrow \text{Min}(v, \text{Max-Value}(s, \text{game}, \alpha, \beta))$ 
    • if  $v \leq \alpha$  then return  $v$  /*  $\alpha$  CUT!! */
    •  $\beta \leftarrow \text{Min}(\beta, v)$ 
end
return  $v$ 

```

Sorry I didn't have a chance to talk about the homework, so
I'll probably send out to beef email with some tips, hey,
zeal, and I again forgot to calculate this

$\alpha - \beta$ Pruning Algorithm: Min-Value



```

function Min-Value (state, game,  $\alpha$ ,  $\beta$ ) return utility value
 $\alpha$ : best MAX on path to state ;  $\beta$ : best MIN on path to state
if Cutoff(state) then return Utility (state)
 $v \leftarrow \infty$ 
for each  $s$  in Successor(state) do
    •  $v \leftarrow \text{Min}(v, \text{Max-Value}(s, \text{game}, \alpha, \beta))$ 
    • if  $v \leq \alpha$  then return  $v$  /*  $\alpha$  CUT!! */
    •  $\beta \leftarrow \text{Min}(\beta, v)$ 
end
return  $v$ 

```

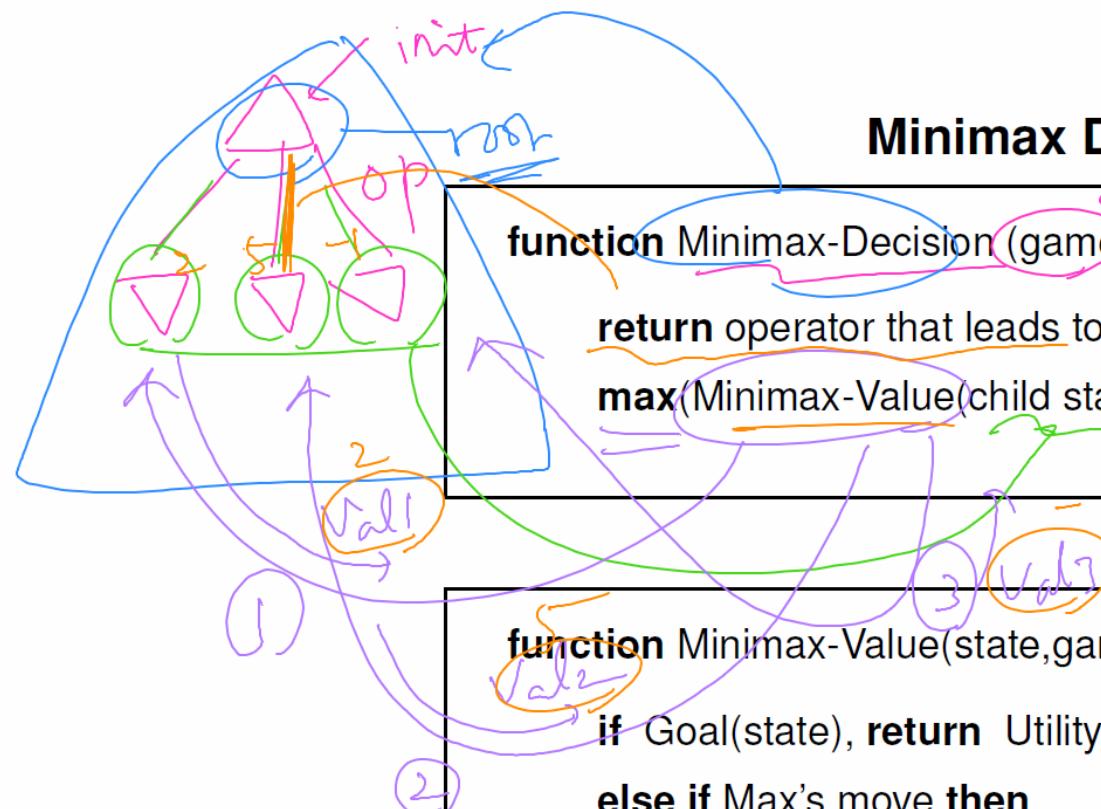
Sorry I didn't have a chance to talk about the homework, so
I'll probably send out to beef email with some tips, hey,
zeal, and I again forgot to calculate this



Minimax Decision

```
function Minimax-Decision(game) returns operator
    return operator that leads to a child state with the
    max(Minimax-Value(child state,game))
```

```
function Minimax-Value(state,game) returns utility value
    if Goal(state), return Utility(state)
    else if Max's move then
        → return max of successors' Minimax-Value
    else
        → return min of successors' Minimax-Value
```



Minimax Decision

```
function Minimax-Decision(game) returns operator
    return operator that leads to a child state with the
    max(Minimax-Value(child state, game))
```

```
function Minimax-Value(state, game) returns utility value
    if Goal(state), return Utility(state)
    else if Max's move then
        → return max of successors' Minimax-Value
    else
        → return min of successors' Minimax-Value
```



Choe, Yoonsuck

Minimax
-dec

Minimax Decision

```
function Minimax-Decision(game) returns operator
    return operator that leads to a child state with the
    max(Minimax-Value(child state, game))
```

```
function Minimax-Value(state, game) returns utility value
    if Goal(state), return Utility(state)
    else if Max's move then
        → return max of successors' Minimax-Value
    else
        → return min of successors' Minimax-Value
```

Then another min. next value on the second child, and so on.



Minimax Decision

function Minimax-Decision (game) **returns** operator

return operator that leads to a child state with the
 max Minimax-Value(child state,game))

function Minimax-Value(state,game) **returns** utility value

if Goal(state), **return** Utility(state)
else if Max's move then
 → **return** max of successors' Minimax-Value
else
 → **return** min of successors' Minimax-Value

